

Safe Region-based Memory Management for Java

Jan Vitek

with Tian Zhao and Jason Baker



(§³)

This talk

- High-level programming languages facilitate software development by abstracting error prone or tedious tasks.
- Memory management is error prone and tedious, but real-time systems sometimes need fine grained control.
- This talk shows how to regain control over memory ... when needed.
- The proposed solution is a hybrid model, a language running mostly on GC but with spurts of manual allocation ... with a new high-level abstraction for Java called Scopes.
- Our solution leverages a long line of results from type theory, the key contribution is a statically safe region-based system which is simple and backwards compatible.
- Some related work: MLKit, Cyclone, Ownership types.

Embedded Real-time Systems

- Real-time embedded systems are central to many applications from avionics, to automotive industry; they are the largest installed base of microprocessors
- The size of embedded systems is growing steadily; up to multi-million line systems (e.g. DD(X) battleship control)
- Very low level languages (e.g. assembly) are not viable; low-level ones (eg. C) are barely tolerable; Ada has unfortunately not found the degree of adoption it deserved...

... new programming language abstractions for embedded real-time systems are sorely needed.

- Reusability is becoming mandatory, even Boeing changing from the old way (recode from scratch) in favor of COTS




Automated Memory Management

- Garbage collection algorithms along with memory-safe languages prevent all memory errors.
- The cost is an increase in memory requirements (or a decrease in performance, take you pick) with 2x over the optimal not uncommon.
- But the main issue is pause times. Stop-the-world collectors require the application to pause for >100ms. In a RT system this can mean missing deadlines.
- Recently, real-time collectors have been able to bound pauses to <20ms, this comes at some cost in throughput
- The state of the art in RT GC seems good enough for many RT applications, but not for some hard real-time subsystems

Region-based allocation

- Region-based allocation is an alternative to GC and manual allocation. It follows from the observation that data exhibits liveness locality (allocation and deallocation times are correlated).
- Idea:** create regions which are pools of memory that can be used to store data and bulk deallocated in constant time.
- This simplifies memory management because instead of tracking individual pointers, we deal with entire regions.
- A stack of regions allows to deal with different lifetimes. Parent regions outlive child regions; data is allocated in the region most closely matching its lifetime.
- In itself, region-based allocation is not safe. Programming language techniques must be devised to enforce safety (either run-time or compile-time).
- The **Real-time Specification for Java** provides region-based MM, but safety is enforced dynamically. This entails runtime overheads and runtime failures.



Motivating example: PRISMj

Mission critical avionics DRE

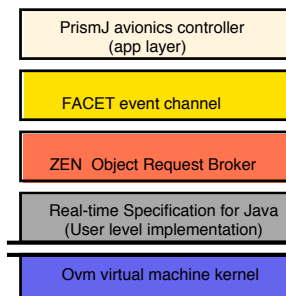
Boeing, Purdue, UCI, WUSTL

Route computation, Threat deconfliction algorithms
ScanEagle UAV

System	K LOCs
PRISMj	109K
FACET EVENT CHANNEL	15K
ZEN CORBA ORB	179K
RTSJ LIBRARIES	60K
CLASSPATH LIBRARIES	500K
OVM VIRTUAL MACHINE	220K



3 rate groups (20, 5, 1 Hz)
several hundred RT threads



kernel
boundary

Embedded Planet PowerPC 8260

Core at 300 MHz
256 Mb SDRAM
32 Mb FLASH
PC/104 mechanical sized
Embedded Linux



SCOPES

Design Requirements

- The design of scopes abides by the following requirements:
 1. **Static type safety:** *no runtime errors and no runtime checks.*
 2. **GC-safety:** *no interference with the GC so that a thread within a scope may safely preempt the garbage collector and never need block for GC.*
 3. **Bytecode compatibility:** *the output of the compiler should be valid Java bytecode. The only allowed change to the tool chain is the addition of an extra verification pass.*
 4. **Minimal changes to the VM:** *modifications should be limited to the memory subsystem.*
 5. **Source-level backwards compatibility:** *existing Java class should be reusable from scope code.*

Scope Basics

- Scope objects reify allocation contexts
- Creating a scope allocates a backing store
- Invoking a method on a scope switch allocation context
- Releasing a scope deallocates the backing store
- nb: use of syntactic sugar can be replaced by annotations

```
1  r = new processor;  
2  while (b = nextRequest()) {  
3      r.getMessage(b,m);  
4      m.dispatch();  
5      release r;  }
```

Programming with Scopes

- Scopes have fields, methods, nested classes and scopes.
- Bound classes are always allocated within their enclosing scope.
- Scopes form a tree at runtime, parents always outlive children.
- Invoking a bound class method switches allocation context.

```
1  scope processor {  
2      class Unpacker { ...  
3          void parse(Bytes b) {...}  
4          void write(Message m) {...}  
5      }  
6      void getMessage(Bytes b, Message m) {  
7          Unpacker p = new Unpacker();  
8          p.parse(b);  
9          p.write(m);  
10 } }
```

A Scope Pool

- Multiple threads can execute in the same scope and communicate via scope fields
- Releasing a scope deallocates the contents only if no thread is active in it

```
1 scope pool {  
2   area[] a;  int cnt;  
3   void create(int i) {  
4     a=new area[i];  
5     for(int j=0;j<i;j++) a[j]=new area;  
6   }  
7   void run(Message m, Bytes data){  
8     int i = cnt++ % a.length;  
9     a[i].run(m, data);  
10    release a[i];  
11  }  
12  scope area {  
13    void run(Message m, Bytes data){  
14      ... // service request  
15  } } }
```


Implicit Polymorphism

- Classes not lexically enclosed within a scope are **free** classes.
- Free classes are *implicitly* scope-polymorphic.
- Instances of free classes can be allocated within any scope, with the restriction that they not escape their enclosing scope.
- Many library classes can be reused.

```
1 class List {  List next;  Object value; }  
2  
3 scope user {  
4   List l;  
5   void create(int i) {  
6     List t;  
7     for(int j=0;j<i;j++)  
8     { t = new List(); t.next =l; l = t; }  
9   } } }
```


Permanent Classes

- Permanent classes --annotated with a **top** qualifier-- exist in immortal memory.
- Instances of a permanent class are never deallocated.
- They are allowed to create top-level scopes.
- Permanent classes can be observed from the heap, but cannot refer to heap objects.
- If a thread is spawned in a permanent object, it is guaranteed not to experience GC-interference.



Crossing Scopes

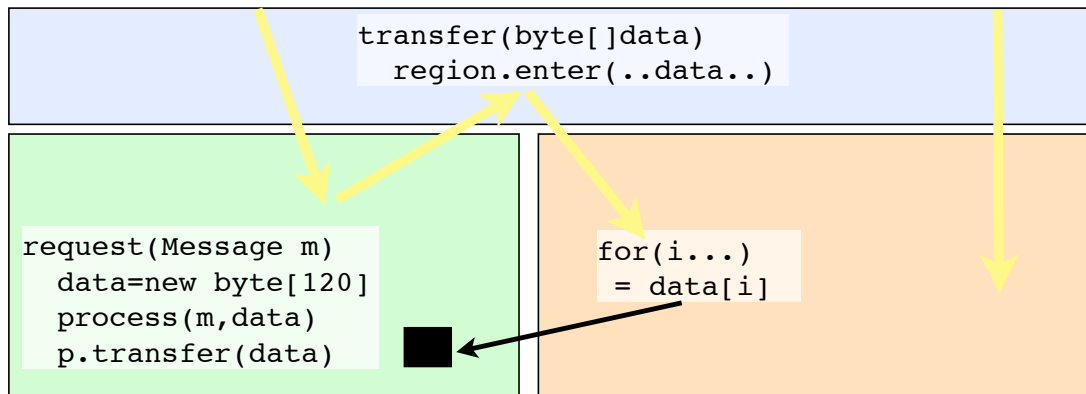
- Bound classes allocated in parent scopes are visible from subscopes.
- Invoking a method of a bound class changes the allocation context to the parent, when the method returns the allocation context is restored.
- Each thread has a stack of scope, but at any given time only one scope is the effective scope.



```

1  scope a {
2      class Box {
3          c s2;
4          void enter() { s2.mc(); }
5      }
6      scope b { void mb(Box box) { box.enter(); } }
7      scope c { void mc(){...} }
8      void ma() {
9          b s1 = new b;  c s2 = new c;
10         Box box = new Box(); box.s2 = s2;
11         s1.mb(box);
12     } }
    
```

Pipelined computations



- LIFO discipline not suited when successive rounds of filtering are needed.
- Pipelined computation can be set up by lending a reference to a sibling region.
- Intuition: a region is pinned while a thread is active within it. If a thread crosses regions, it can safely access its origin scope.
- Safe as long as sibling references are not retained.

Quasi Linear Types

```

1  scope r {
2      class Bridge { q qscp;
3          void run(p pscp, q qscp) {
4              this.qscp=qscp; pscp.enter(this);
5          }
6          void handoff(borrow byte[] data) {
7              qscp.enter( data);
8          }
9      }
10     scope p {
11         void enter(Bridge b) {
12             byte[] data=new byte[20];
13             ... // some computation
14             b.handoff( data);
15         }
16     }
17     scope q {
18         void enter(borrow byte[] data) {
19             ... //use cross-scope reference
20         }
21     }
22 }

```


Quasi Linear Types


- **borrow** -- is a type qualifier that can be used on method arguments and locals.
- The type system ensures a borrowed reference doesn't outlive the method invocation.
- It is safe to access a borrowed reference from any scope.
- Borrowed references cannot be assigned to fields of an object or scope.
- It's safe to modify primitive fields of a borrowed object, but not reference fields.
- Any reference retrieved from a borrowed reference is borrowed as well.

```
class BL{BL next;int i;F f;}

method_m( borrow BL bl ) {
    borrow nbl = bl.next;    // ok
    bl.i++;                  // ok
    bl.f = new F();          // Unsafe! Compile time error.
```


GC-safety

- In a hybrid setting, hard RT code mustn't pause for GC.
- Two scenarios for GC pauses
 - ☑ real-time code allocates in the heap
 - ☑ real-time thread released in midst of GC:
wait for GC to finish or observe inconsistent data
- Scopes can not refer to the heap thus neither can happen
- Do not prevent transitive priority inversion scenarios
 - 1) a non-rt thread grabs a lock on a permanent object,
 - 2) GC is triggered,
 - 3) a RT thread tries to access the locked object.



The Type System

- The type system enforces the following invariants:



- (1) *Scopes* ($s, s.p, s.q$) may refer to objects within their allocation context and to bound classes of ancestor scopes.
- (2) *Bound classes* (e.g. C, D) may refer to any co-located object, and to bound classes allocated in a parent scope.
- (3) *Free class instances allocated in a scope* (e.g. B) may refer to any co-located object.
- (4) *Free class instances allocated in the heap* (e.g. A) may refer to permanent classes.
- (5) *Permanent classes* (e.g. E) may refer to other permanent classes and top-level scopes.
- (6) *Borrowed reference* may refer to any bound class.

The Type System

- Some of the constraints are:
 - ☑ Scopes can not extend other classes
 - ☑ Bound classes can only extend bound classes in the same scope (or free classes)
 - ☑ Free classes can only extend free classes.
 - ☑ Borrowed types can not be stored in fields, or have reference types stored into.
 - ☑ Widening can not cast off the borrowed annotation.
 - ☑ Reference types returned by a borrowed object must be borrowed.
 - ☑ Free types are globally visible
 - ☑ Bound types are visible in the defining scope and subscopes
 - ☑ Scope types are only visible in the enclosing scope
 - ☑ The public interface (field/methods) of scopes and bound types can not contain free classes
 - ☑ Widening is restricted in a number of ways ...

Proof of Soundness

- The scope type system has been formalized as an extension of Igarashi e.a.'s Featherweight Java, and proven sound. The main theorem ensures that after a release, there cannot be any reference into the scope.

$$\begin{array}{ll}
 R ::= \text{scope } S \{ \overline{K} \overline{f}; \overline{M} \} & \frac{\Gamma \vdash_{K_t} e : T \quad \text{fields}(T) = (\overline{K} \overline{f}) \quad T \upharpoonright_{K_t} = T' \quad (T, K_t) <: (K_t, T') \quad T' \text{ viz } K_t}{\Gamma \vdash_{K_t} e.f_1 : T'} \quad (\text{T-FIELD}) \\
 L ::= \text{class } C \triangleleft C' \{ \overline{K} \overline{f}; \overline{M} \} & \frac{\Gamma \vdash_{K_t} e : K \quad \text{fields}(K) = (\overline{K} \overline{f}) \quad \Gamma \vdash_{K_t} e' : K' \quad (K_t, K') <: (K_t, K)}{\Gamma, \Sigma \vdash_{K_t} e.f_1 = e' : K'} \quad (\text{T-UPD}) \\
 M ::= K \text{ m } (\overline{T} \overline{x}) \{ \text{return } e; \} & \\
 e ::= x \mid v \mid \text{new } K() \mid e.f \mid e.f := e \mid e.m(\overline{e}) & \frac{K \neq c \Rightarrow (\text{scopeof}(K_t) = S \quad K = S.c \vee K = S.s)}{\Gamma \vdash_{K_t} \text{new } K() : K} \quad (\text{T-NEW}) \\
 & \mid (K) e \mid \text{release } e \\
 v ::= \ell \mid \text{null} & \frac{\Gamma \vdash_{K_t} e : T_r \quad \Gamma \vdash_{K_t} \overline{e} : \overline{T}' \quad \text{mtype}(m, T_r) = \overline{T}' \rightarrow K \quad T_r \upharpoonright K = T \quad T \text{ viz } K_t}{\forall i, (K_t, T'_i) <: (T_r, T_i) \quad (T_r, K) <: (K_t, T)} \quad (\text{T-INVK}) \\
 T ::= K \mid \text{borrow } C & \frac{\Gamma \vdash_{K_t} e : K' \quad (K = c \Rightarrow K' = c' \vee \text{scopeof}(K') = \text{scopeof}(K_t))}{\Gamma \vdash_{K_t} (K) e : K} \quad (\text{T-CAST}) \\
 K ::= C \mid S & \\
 C ::= S.c \mid c & \\
 S ::= S.s \mid \text{top} & \frac{\Gamma \vdash_{K_t} x : \Gamma(x)}{\Gamma \vdash_{K_t} x : \Gamma(x)} \quad (\text{T-VAR}) \quad \frac{\Gamma \vdash_{K_t} e : S}{\Gamma \vdash_{K_t} \text{release } e : S} \quad (\text{T-REL})
 \end{array}$$

Method typing:

$$\frac{\overline{x} : \overline{T}, \text{this} : C, \emptyset \vdash_C e : K' \quad K, \overline{T} \text{ viz } C \quad \text{override}(m, C, C') \quad (C, K') <: (C, K)}{C \triangleleft C' \vdash K \text{ m } (\overline{T} \overline{x}) \{ \text{return } e; \}} \quad (\text{T-METHC}) \quad \frac{\overline{x} : \overline{T}, \text{this} : S, \emptyset \vdash_S e : K' \quad K, \overline{T} \text{ viz } S \quad (S, K') <: (S, K)}{S \vdash K \text{ m } (\overline{T} \overline{x}) \{ \text{return } e; \}} \quad (\text{T-METHS})$$

Class typing:

$$\frac{C \triangleleft C' \vdash \overline{M} \quad \overline{K} \text{ viz } C \quad C' = c \vee (C = S.c \wedge C' = S.c')}{\text{class } C \triangleleft C' \{ \overline{K} \overline{f}; \overline{M} \} \text{ OK}} \quad (\text{T-CLASS}) \quad \frac{S \vdash \overline{M} \quad \overline{K} \text{ viz } S}{\text{scope } S \{ \overline{K} \overline{f}; \overline{M} \} \text{ OK}} \quad (\text{T-SCOPE})$$

Conclusions

- High-level languages for real-time systems must support memory-safe, and GC-safe, manual memory management techniques.
- Region-based allocation provides linear time allocation and bulk deallocation.
- Type systems can prevent all memory violations at compile time.
- The scopes abstraction introduced in this talk provides static safety without requiring drastic changes to Java.
- Scopes are simple to use and allow the reuse of existing code.