

# Lecture: Embedded Software Architectures

Jan Vitek

ECE/CS  
Spring 2011

## Reading List

2

- Mandatory Reading
  - Chapter 5 of ECS textbook
- Optional Reading
  - N/A

## Software Architecture

3

- A *software architecture* gives the general structure of an embedded application independent of the actual computation performed
- Choice of architecture impacts issues such as:
  - development time / likelihood of software defects
  - responsiveness and latency
  - code size / complexity
- Rule of thumb:
  - *Select simplest architecture that meets application requirements*
  - Any extraneous complexity/generalality costs additional development and verification effort

# Software Architectures

4

- Four well known choices:
  - Simple Round Robin
  - Round Robin with Interrupts
  - Round Robin with Interrupts and Function Queues
  - Real-time Operating System-based architectures
- The architectures are sorted in order of increasing complexity
- Round Robin (RR) architectures are also called *Cyclic Executives* in real-time literature
- The main different between RR and RTOS-based approaches is that in RR scheduling and admission control is done by the developer as opposed to leaving it to the OS

## Round Robin

5

- Simplest architecture, a single loop checks devices in predefined sequence and performs I/O right away

```
1. while(1) {
2.   if (device_1_ready()) { /*Perform D1 I/O and relate computation.*/ }
3.   if (device_2_ready()) { /*Perform D2 I/O and relate computation.*/ }
4.   ...
5.   if (device_N_ready()) { /*Perform DN I/O and relate computation.*/ }
6. }
```

- Works well for system with few devices, trivial timing constraints, proportionally small processing costs
- Response time of device  $i$  equal to WCET of the body of the loop

## Round Robin

6

- Periodic Round Robin
  - In case the system must perform operations at different frequencies
  - Add code to wait a variable amount of time

```
1.while(1) {
2.  waitForNextPeriod(10); // idle for up to 10 ms
3.  if (device_1_ready()) { /*Perform D1 I/O and relate computation.*/ }
4.  ...
}
```

- Exercise:
  - Think of how to implement a loop that runs every 10 ms and measures the drift

# Round Robin

- Limitations of the architecture:
  - If some devices require small response times, while other have large WCET it will not be possible to guarantee that all timing constraints will be met
  - The architecture is fragile, adding a new task can easily cause missed deadlines

- Question:
  - Is the order in which devices appear significant?
  - Same question, but with code for devices having different processing times and timing constraints?


# Round Robin with Interrupts

- Hardware events requiring small response times handled by ISRs
- Typically ISRs do little more than set flags and copy data

```
1. bool f_device_1 = FALSE;
2. bool f_device_2 = FALSE;
3. ...
4. void interrupt handle_dev_1() {
5.     // handle device 1
6.     f_device_1 = TRUE;
7. }
8. ...
9. void main() {
10.    while (1) {
11.        if(f_device_1) {
12.            f_device_1 = FALSE;
13.            // do processing related to device 1...
14.            if (f_device_2) {
15.                ...
16.            }
17.    }
```

# Round Robin with Interrupts

- Latency of an ISR is function of response time of higher priority ISRs
- Lower bound on latency of RR loop is response time of the ISRs



## Round Robin with Interrupts

10

- Drawbacks

- All task code executes at same priority
  - One can test some flags multiples times within loop body to reduce latency
- Shared data bugs

- Question:

- What if one of the device requires large amount of processing time (larger than the time constraint of others?)

## RR+I and Function Queues

11

- Rather than fixed order, program manages order of execution

```
1. #define DEV_1 1
2. #define DEV_2 2
3. ...
4. void interrupt handle_DEV_1() {
5.     // deal with device
6.     enqueue(DEV_1);
7. }
8. ...
9. void main() {
10.    while (1) {
11.        switch (dequeue()) {
12.            case DEV_1:    // process DEV_1
13.                break;
14.            ...
15.            default:      // empty queue nothing to do
16.        }
17.    }
18.}
```

## RR+I and Function Queues

12

- One could use function pointers, but they add complexity
  - FP are useful if one does not want to hardwire the devices in the main loop
- enqueue() reorders queue to improve latency of high priority devices
- For long running functions: break them up into multiple smaller units
  - Question: does that improve latency?

- Question

- Consider implementation of dequeue(), what kind of data structure would you use (why), is special care needed?

# Real-time Operating Systems

12

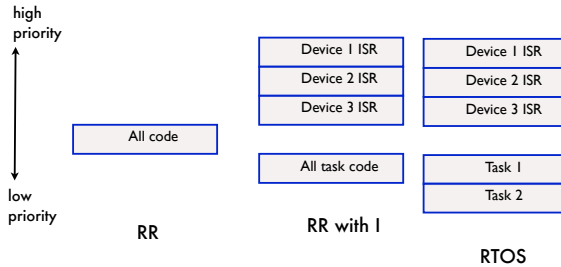
- Rely on the operating for scheduling tasks
- Leverage preemptive scheduling to ensure that deadlines are met

```
1. static pthread_t thread_1;  
2. ...  
3. void interrupt handle_DEV_1() {  
4.     // handle device  
5.     CHECK( pthread_wakeup(thread_1) );  
6. }  
7. ...  
8. void task_1() {  
9.     while (1) {  
10.        pthread_suspend_np();  
11.        // process device 1 I/O  
12.    }  
13. }  
14. ...
```

# Real-time Operating Systems

14

- The scheduler in a RTOS takes care of scheduling all tasks according to their priority
- Long running, low priority, tasks can be preempted by higher priority ones



# Real-time Operating Systems

15

- Services that an RTOS *could* provide:

- Scheduling tasks
  - create/terminate threads
  - timing threads operations
  - preemption
- Synchronization
  - semaphores and locks
- Input/Output
  - interrupt handling
- Memory management
  - separate stacks
  - segmentation
  - allocation/deallocation
- File system
  - persistent store
- Security
  - user vs. kernel space
  - identity management

# Conclusion

16

- Software architectures describe the structure of a system independently of its function
  - *Round Robin* is a simple architecture for devices with few (or uniform) timing constraints
  - *Round Robin with Interrupts* extends RR with low-latency interrupt handling
  - *Round Robin with Interrupts and Function Queues* allows dynamic scheduling of tasks under programmatic control
  - *Real-time Operating Systems* relieve programmers from having to deal with scheduling and time management