

# **The Java Virtual machine**

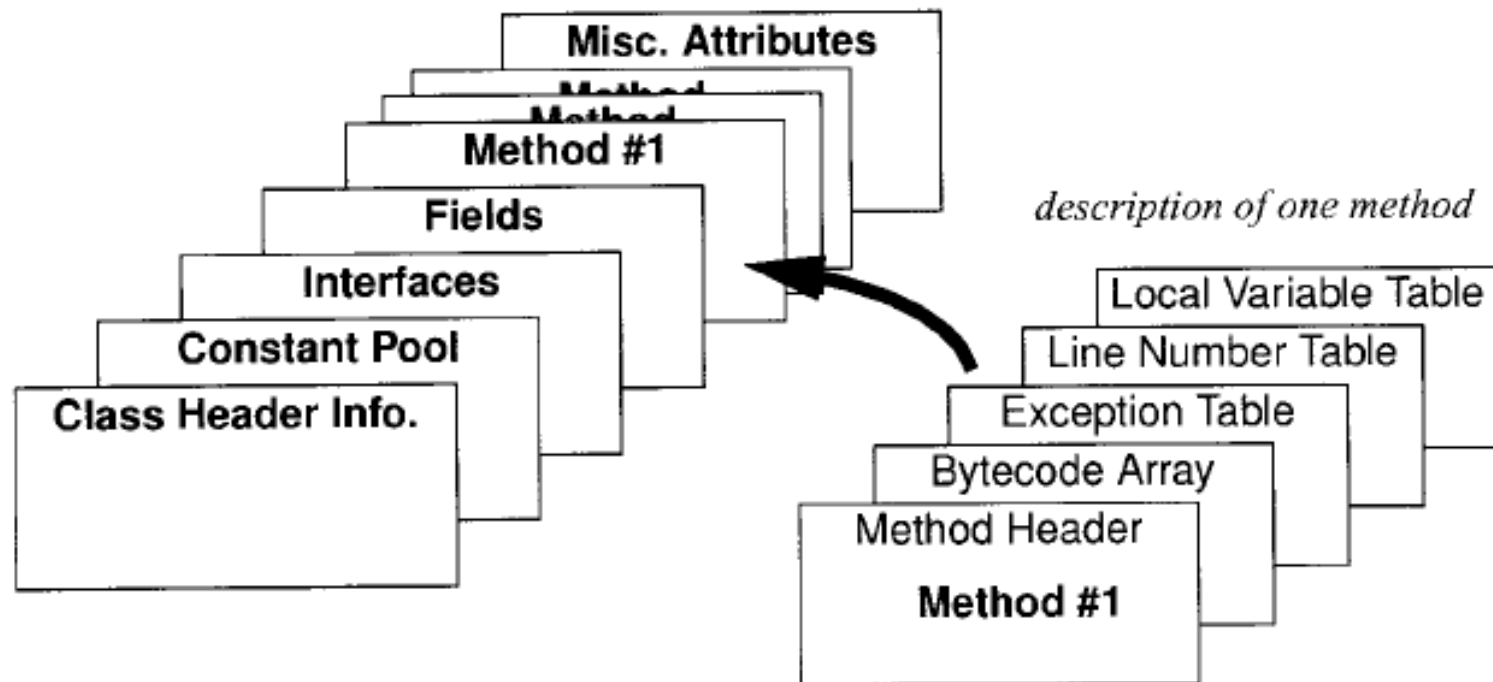
## **A Main Reference Source**

**The Java™ Virtual Machine Specification (2nd Ed)  
by Tim Lindholm & Frank Yellin  
Addison-Wesley, 1999**

**The book is on-line and available for download:**

**<http://java.sun.com/docs/books/vmspec/>**

# The Java Classfile



## JVM Runtime Behaviour

- VM startup
- Class Loading/Linking/Initialization
- Instance Creation/Finalisation
- Unloading Classes
- VM exit

# VM Startup and Exit

## Startup

- Load, link, initialize class containing `main()`
- Invoke `main()` passing it the command-line arguments
- Exit when:
  - all non-daemon threads end, or
  - some thread explicitly calls the `exit()` method

## Class Loading

- Find the binary code for a class and create a corresponding Class object
- Done by a class loader – bootstrap, or create your own
- Optimize: prefetching, group loading, caching
- Each class-loader maintains its own namespace
- Errors include: `ClassFormatError`, `UnsupportedClassVersionError`, `ClassCircularityError`, `NoClassDefFoundError`

## Class Loaders

- System classes are automatically loaded by the bootstrap class loader
- To see which:  
`java -verbose:class Test.java`
- Arrays are created by the VM, not by a class loader
- A class is unloaded when its class loader becomes unreachable (the bootstrap class loader is never unreachable)

## Class Linking - 1. Verification

- Extensive checks that the .classfile is valid
- This is a vital part of the JVM security model
- Needed because of possibility of:
  - buggy compiler, or no compiler at all
  - malicious intent
  - (class) version skew
- Checks are independent of compiler and language



## Class Linking - 2. Preparation

- Create static fields for a class
- Set these fields to the standard default values (N.B. not explicit initializers)
- Construct method tables for a class
- ... and anything else that might improve efficiency

## Class Linking - 3. Resolution

- Most classes refer to methods/fields from other classes
- Resolution translates these names into explicit references
- Also checks for field/method existence and whether access is allowed

## Class Initialization

Happens *once* just before first instance creation, or first use of static variable.

- Initialise the superclass first!
- Execute (class) static initializer code
- Execute explicit initializers for static variables
- May not need to happen for use of *final* static variable
- Completed before anything else sees this class

## Instance Creation/Finalisation

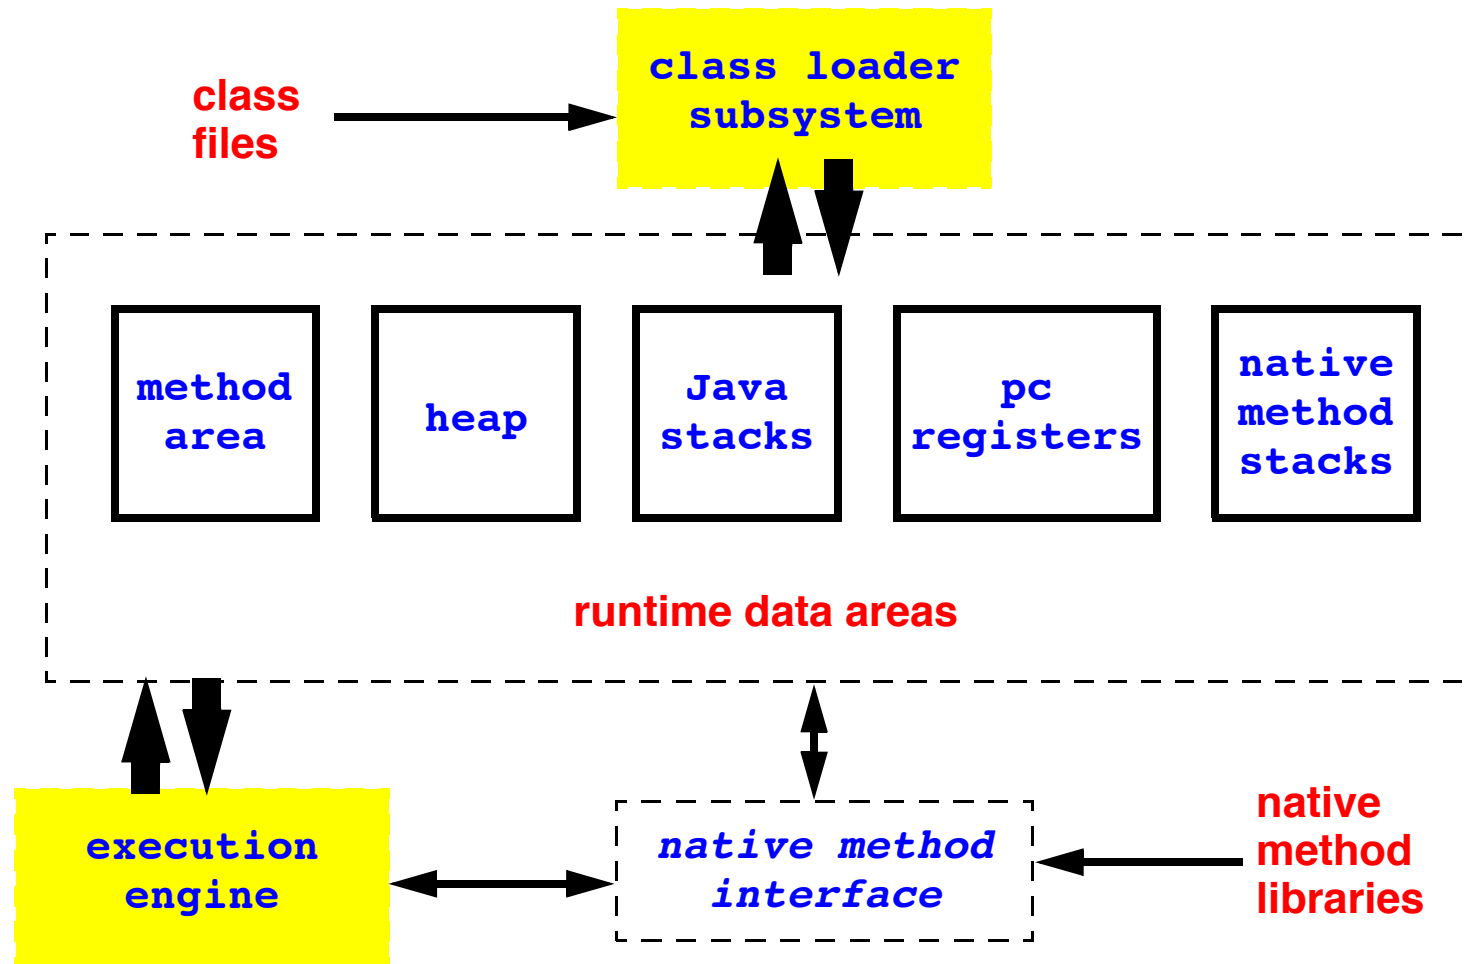
- Instances are created using `new`, or `newInstance()` from class `Class`
- Instances of `String` may be created (implicitly) for String literals
- Process:
  - 1 Allocate space for all the instance variables (including the inherited ones),
  - 2 Initialize them with the default values
  - 3 Call the appropriate constructor (do parent's first)
- `__finalize()` is called just before garbage collector takes the object (so timing is unpredictable)

## JVM Architecture

**The internal runtime structure of the JVM consists of:**

- One: (i.e. shared by all threads)
  - method area
  - heap
- For each thread, a:
  - program counter (pointing into the method area)
  - Java stack
  - native method stack (system dependent)

## Run-Time Data Areas (Venners Figure 5-1)



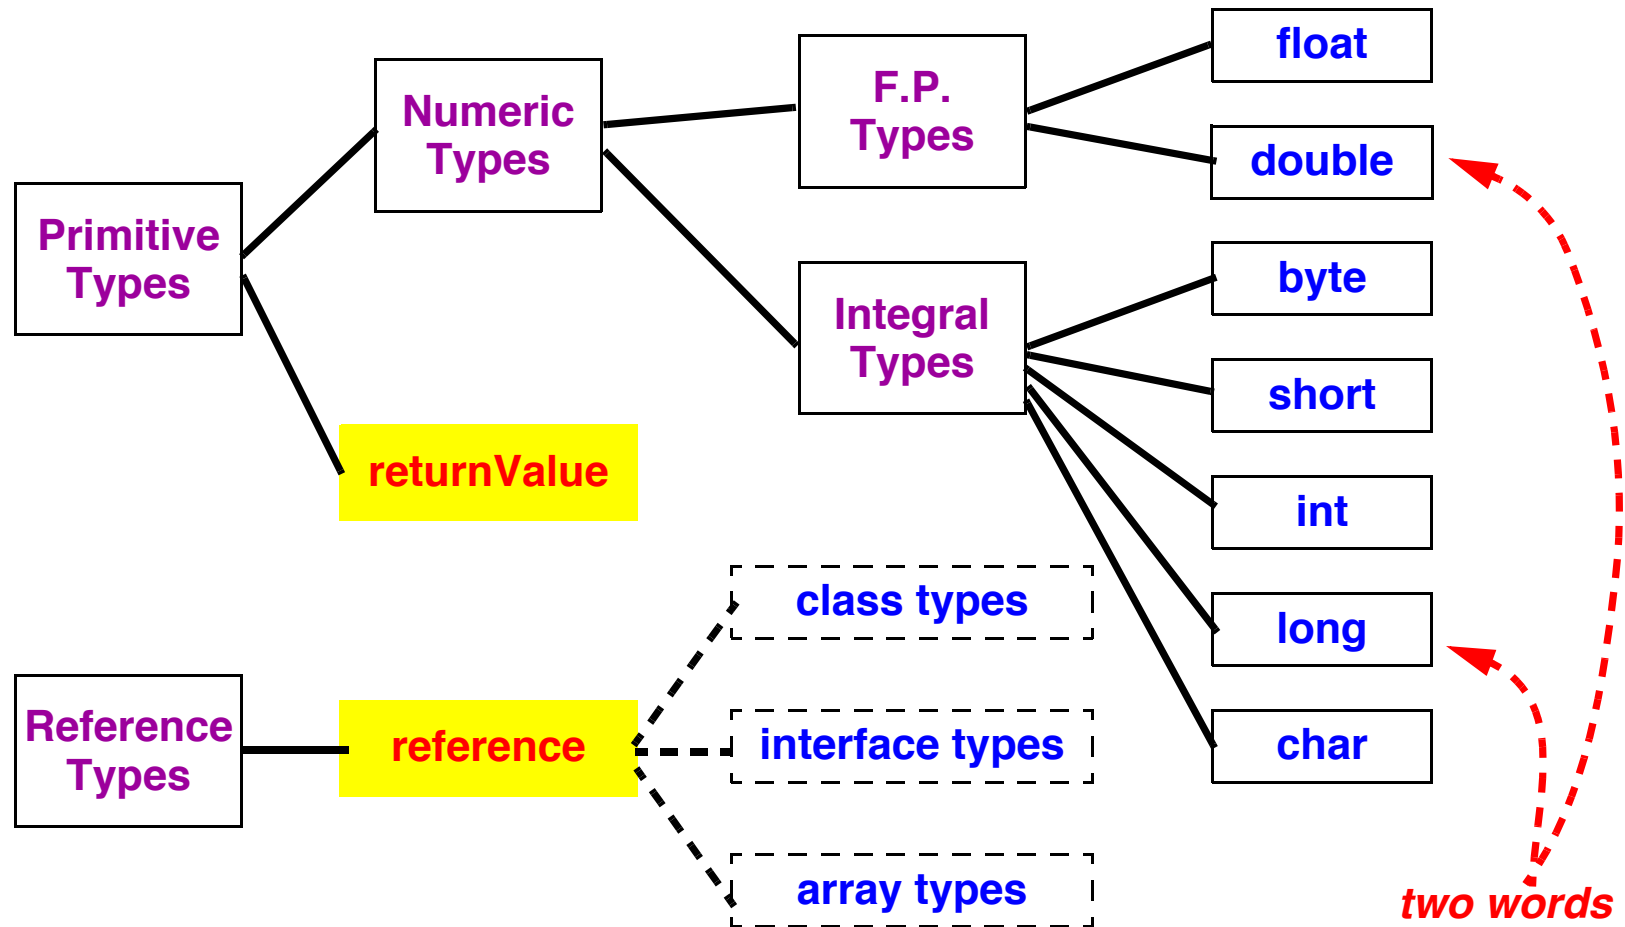
# Java Bytecode

# Java Intermediate Bytecode

- By James Gosling; presented at IR'95.
- Quick overview:
  - argue for the presence of type information in the bytecode
  - benefits for checkability (because speed/security)
  - reduced dependencies on environment



# Datatypes of the JVM (Venners 5-4)



## Control Transfer

- ifeq, iflt, ifle, ifne, ifgt, ifge
- ifnull, ifnonnull
- if\_icmpeq, if\_icmplt, if\_icmple, if\_icmpne, if\_icmpgt, if\_icmpge
- if\_acmpeq, if\_acmpne
- goto, goto\_w, jsr, jsr\_w, ret

## Switch statement implementation

- tableswitch, lookupswitch

## Comparison operations for long, float & double types

- lcmp, fcmpl, fcmpg, dcmpl, dcmpg

# Load and Store Instructions

## Transferring values between local variables and operand stack

- iload, lload, fload, dload, aload  
and special cases of the above: iload\_0, iload\_1 ...
- istore, lstore, fstore, dstore, astore

## Pushing constants onto the operand stack

- bipush, sipush, ldc, ldc\_w, ldc2\_w, aconst\_null, iconst\_m1  
and special cases: iconst\_0, iconst\_1, ...

## Arithmetic Operations

**Operands are normally taken from operand stack and the result pushed back there**

- iadd, ladd, fadd, dadd
- isub ...
- imul ...
- idiv ...
- irem ...
- ineg ...
- iinc

## Bitwise Operations

- ior, lor
- iand, land
- ixor, lxor
- ishl, lshl
- ishr, iushr, lshr, lushr

## Type Conversion Operations

### Widening Operations

- i2l, i2f, i2d, l2f, l2d, f2d

### Narrowing Operations

- i2b, i2c, i2s, l2i, f2i, f2l, d2i, d2l, d2f

## Operand Stack Management

- pop, pop2
- dup, dup2, dup\_x1, dup\_x2, dup2\_x2, swap

## Object Creation and manipulation

- new
- newarray, anewarray, multinewarray
- getfield, putfield, getstatic, putstatic
- baload, caload, saload, iaload, laload, faload, daload, aaload
- bastore, castore, sastore, iastore, lastore, fastore, dastore, aastore
- arraylength
- instanceof, checkcast

## Object Creation and manipulation

- new
- newarray, anewarray, multinewarray
- getfield, putfield, getstatic, putstatic
- baload, caload, saload, iaload, laload, faload, daload, aaload
- bastore, castore, sastore, iastore, lastore, fastore, dastore, aastore
- arraylength
- instanceof, checkcast

## Method Invocation / Return

- invokevirtual
- invokespecial
- invokeinterface
- invokestatic
- ireturn, freturn, dreturn, areturn
- return

# Java Intermediate Bytecode

- Observation:
  - Original goals where modularity, small footprint, verifiability, but not speed.
  - the bytecode had to be statically typed (speed/safety argument)
  - control flow merges must have the same incoming stack types
  - use symbolic references to environment (fragile base class)



# Class Resolution

- CP entry tagged `CONSTANT_Class` can be either class/interface.
- Execution of an instruction that refers to a class:
  1. *search for class in the classloader hierarchy*
  2. *if not found, initiate class loading*
- ... much more to the story.

# Method Invocation

INVOKEVIRTUAL,	- instance method
INVOKEINTERFACE,	- interface method
INVOKESPECIAL	- constructor/private/super method
INVOKESTATIC	- class method

```
foo/baz/Myclass/myMethod(Ljava/lang/String;)V
```

-----

class          method          descriptor

- When an invocation is executed the method must be resolved.

# Method Resolution

1. Checks if C is class or interface.

If C is interface, throw `IncompatibleClassChangeError`.

2. Look up the referenced method in C and superclasses:

- Success if C has method with same name & descriptor
- Otherwise, if C has a superclass, repeat 2 on super.

3. Otherwise, locate method in a superinterface of C

- If found success.
- Otherwise, fail.

# Method Invocation

- Resolution is rather work intensive. Can this be done faster?

# class initialization

- Before use of static field, static method, object creation, a class must be initialized.
- Initialization involves creating a new Class object, and running the static initializers.
- Every operation that could trigger initialization must check the status of the class.

# subroutines

- Subroutines were added to the bytecode to reduce the space requirements of exception handler's finally clauses.

# Example

```
int bar(int i) {
    try {
        if (i == 3) return this.foo();
    } finally {
        this.ladida();
    }
    return i;
}
```

Region	Target
1–12	17
13–16	21

```
01  iload_1           // Push i
02  iconst_3         // Push 3
03  if_icmpne 10      // Goto 10 if i does not e

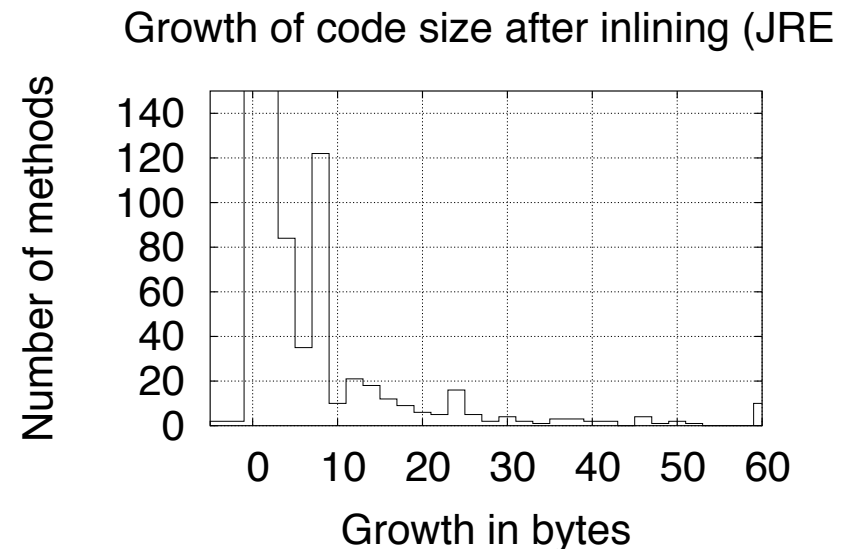

04  aload_0          // Push this
05  invokevirtual foo // Call this.foo
06  istore_2         // Save result of this.foo
07  jsr 13           // Do finally block before
08  iload_2          // Recall result from this
09  ireturn          // Return result of this.f

10  jsr 13           // Do finally block before

11  iload_1          // Push i
12  ireturn          // Return i
    // finally block
13  astore_3         // Save return address in
14  aload_0          // Push this
15  invokevirtual ladida // Call this.ladida()
16  ret 3            // Return to address saved
    // Exception handler for try body
17  astore_2         // Save exception
18  jsr 13           // Do finally block
19  aload_2          // Recall exception
20  athrow           // Rethrow exception
    // Exception handler for finally body
21  athrow           // Rethrow exception
```

# subroutines

- Over the JDK 1.1, subroutines save a total of 2427 bytes [Freund98].
- Java5 does not use them. They can be inlined by tools.



From Artho, Biere, Bytecode 2005.

Figure 7. Sizes of subroutines and size increase after inlining.



# class compression

- Observation:
  - class file size dominated by symbolic information in the CP
  - JAR files (containing multiple classes) contain redundancies

	swingall	javac	[Pugh99]
Total size	3,265	516	
excluding jar overhead	3,010	485	
Field definitions	36	7	
Method definitions	97	10	
Code	768	114	
Other	72	12	
Constant pool	2,037	342	
Utf8 entries	1,704	295	
if shared	372	56	
if shared and factored	235	26	

# compression

- Observation:
  - class file size dominated by symbolic information in the CP
  - JAR files (containing multiple classes) contain redundancies

File Format	Size	% orig. size
JAR file, uncompressed	260,178	100.0%
JAR file, compressed	132,600	51.0%
Clazz	97,341	37.4%
Gzip	97,223	37.4%
Jazz	59,321	22.8%

[Bradley,Horspool,Vitek,98]

## **Java Virtual Machine, part three**

## Verification

- Ensures that the type (i.e. the loaded class) obeys Java semantics, and
- will not violate the integrity of the JVM.

**There are many aspects to verification**

## Verification, cont'd

### Some Checks during Loading

- If it's a classfile, check the magic number (**0xCAFEBAFE**),
- make sure that the file parses into its components correctly

### Additional Checks after/during Loading

- make sure the class has a superclass (only Object does not)
- make sure the superclass is not final
- make sure final methods are not overridden
- if a nonabstract class, make sure all methods are implemented
- make sure there are no incompatible methods
- make sure constant pool entries are consistent

## Additional Checks after/during Loading, cont'd

- check the format of special strings in the constant pool (such as method signatures etc)

## A Final Check (required before method is executed)

- verify the integrity of the method's bytecode

**This last check is very complicated (so complicated that Sun got it wrong a few times)**

# Verifying Bytecode

## The requirements

- All the opcodes are valid, all operands (e.g. number of a field or a local variable) are in range.
- Every control transfer operation (goto, ifne, ...) must have a destination which is in range and is the start of an instruction
- Type correctness: every operation receives operands with the correct datatypes
- No stack overflow or underflow
- A local variable can never be used before it has been initialized
- Object initialization – the constructor must be invoked before the class instance is used

## The requirements, cont'd

- Execution cannot fall off the end of the code
- The code does not end in the middle of an instruction
- For each exception handler, the start and end points must be at the beginnings of instructions, and the start must be before the end
- Exception handler code must start at the beginning of an instruction



## Sun's Verification Algorithm

**A *before* state is associated with each instruction.**

**The state is:**

- contents of operand stack (stack height, and datatype of each element), plus
- contents of local variables (for each variable, we record *uninitialized* or *unusable* or the datatype)

**A datatype is integral, long, float, double or any reference type**

**Each instruction has an associated *changed* bit:**

- all these bits are false,
- except the first instruction whose changed bit is true.

## Sun's Verification Algorithm, cont'd

```
do forever {  
  find an instruction I whose changed bit is true;  
  if no such instruction exists, return SUCCESS;  
  set changed bit of I to false;  
  state S = before state of I;  
  
  for each operand on stack used by I  
    verify that the stack element in S has correct datatype  
    and pop the datatype from the stack in S;  
  for each local variable used by I  
    verify that the variable is initialized and  
    has the correct datatype in S;  
  if I pushes a result on the stack,  
    verify that the stack in S does not overflow, and  
    push the datatype onto the stack in S;  
  if I modifies a local variable,  
    record the datatype of the variable in S  
    ... continued
```

## Sun's Verification Algorithm, cont'd

determine SUCC, the set of instructions which can follow I;  
(Note: this includes exception handlers for I)

for each instruction J in SUCC do

merge *next* state of I with the *before* state of J

and set J's *changed* bit if the *before* state changed;

(Special case: if J is a destination because of an exception  
then a special stack state containing a single instance of  
the exception object is created for merging with the *before*  
state of J.)

} // end of do forever

**Verification fails if a datatype does not match with what is required by the instruction, the stack underflows or overflows, or if two states cannot be merged because the two stacks have different heights.**

## Sun's Verification Algorithm, cont'd

### Merging two states

- Two stack states with the same height are merged by pairwise merging the types of corresponding elements.
- The states of the two sets of local variables are merged by merging the types of corresponding variables.

### The result of merging two types:

- Two types which are identical merge to give the same type
- For two types which are not identical:  
if they are both references, then the result is the first common superclass (lowest common ancestor in class hierarchy);  
otherwise the result is recorded as *unusable*.

## Example (Leroy, Figure 1):

```
static int factorial( int n ) {
    int res;
    for (res = 1; n > 0; n--) res = res * n;
    return res;
}
```

## Corresponding JVM bytecode:

```
method static int factorial(int), 2 variables, 2 stack slots
  0: iconst_1          // push the integer constant 1
  1: istore_1          // store it in variable 1 (res)
  2: iload_0           // push variable 0 (the n parameter)
  3: ifle 14           // if negative or null, go to PC 14
  6: iload_1           // push variable 1 (res)
  7: iload_0           // push variable 0 (n)
  8: imul              // multiply the two integers at top of stack
  9: istore_1          // pop result and store it in variable 1
10: iinc 0, -1         // decrement variable 0 (n) by 1
11: goto 2            // go to PC 2
14: iload_1           // load variable 1 (res)
15: ireturn            // return its value to caller
```

# Sun's Analysis Algorithm

Chng'd	State before		Instruction	State after	
	Stack	Locals		Stack	Locals
X	()	(I,T)	0: <code>iconst_1</code>		
-	?	(?,?)	1: <code>istore_1</code>		
-	?	(?,?)	2: <code>iload_0</code>		
-	?	(?,?)	3: <code>ifl 14</code>		
-	?	(?,?)	6: <code>iload_1</code>		
-	?	(?,?)	7: <code>iload_0</code>		
-	?	(?,?)	8: <code>imul</code>		
-	?	(?,?)	9: <code>istore_1</code>		
-	?	(?,?)	10: <code>iinc 0, -1</code>		
-	?	(?,?)	11: <code>goto 2</code>		
-	?	(?,?)	14: <code>iload_1</code>		
-	?	(?,?)	15: <code>ireturn</code>		

where **I** = integral; **T** = *uninitialized/unusable*; ? = **⊥** = *unknown*

# Sun's Analysis Algorithm - after 1 step

Chng'd	State before		Instruction	State after	
	Stack	Locals		Stack	Locals
-	()	(I,T)	0: <code>iconst_1</code>	(I)	(I,T)
X	(I)	(I,T)	1: <code>istore_1</code>		
-	?	(?,?)	2: <code>iload_0</code>		
-	?	(?,?)	3: <code>iflt 14</code>		
-	?	(?,?)	6: <code>iload_1</code>		
-	?	(?,?)	7: <code>iload_0</code>		
-	?	(?,?)	8: <code>imul</code>		
-	?	(?,?)	9: <code>istore_1</code>		
-	?	(?,?)	10: <code>iinc 0, -1</code>		
-	?	(?,?)	11: <code>goto 2</code>		
-	?	(?,?)	14: <code>iload_1</code>		
-	?	(?,?)	15: <code>ireturn</code>		

# Sun's Analysis Algorithm - after 4 steps

Chng'd	State before		Instruction	State after	
	Stack	Locals		Stack	Locals
-	()	(I,T)	0: <code>iconst_1</code>		
-	(I)	(I,T)	1: <code>istore_1</code>		
-	()	(I,I)	2: <code>iload_0</code>		
-	(I)	(I,I)	3: <code>ifle 14</code>	(I)	(I,I)
X	()	(I,I)	6: <code>iload_1</code>		
-	?	(?,?)	7: <code>iload_0</code>		
-	?	(?,?)	8: <code>imul</code>		
-	?	(?,?)	9: <code>istore_1</code>		
-	?	(?,?)	10: <code>iinc 0, -1</code>		
-	?	(?,?)	11: <code>goto 2</code>		
X	(I)	(I,I)	14: <code>iload_1</code>		
-	?	(?,?)	15: <code>ireturn</code>		

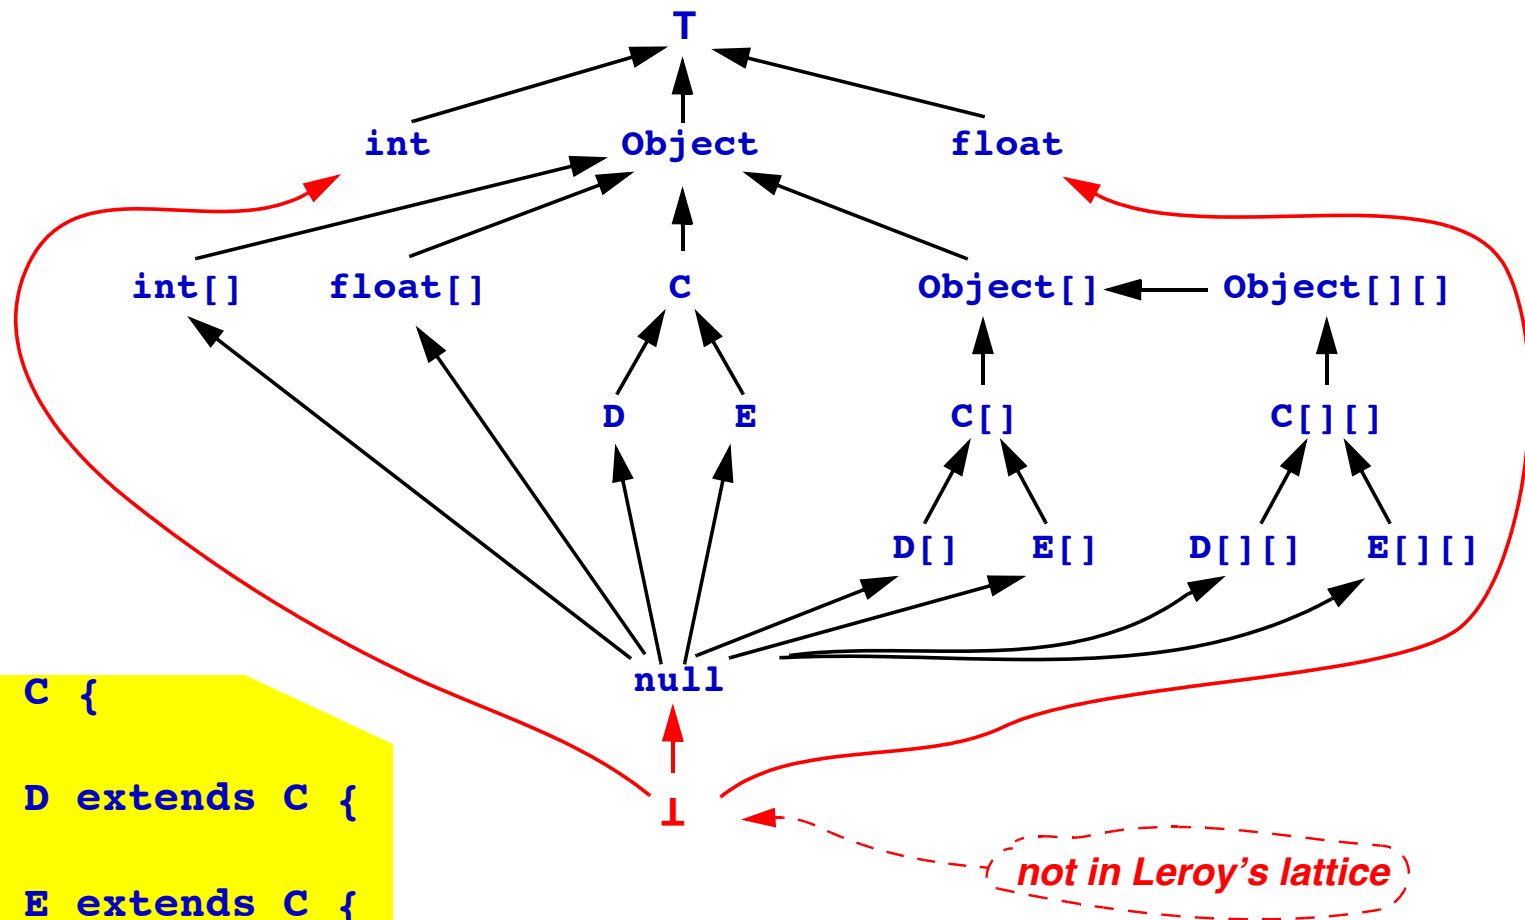


## Analysis Algorithm - after 12 steps

Chng'd	State before		Instruction	State after	
	Stack	Locals		Stack	Locals
-	()	(I,T)	0: <code>iconst_1</code>		
-	(I)	(I,T)	1: <code>istore_1</code>		
-	()	(I,I)	2: <code>iload_0</code>		
-	(I)	(I,I)	3: <code>ifle 14</code>		
-	()	(I,I)	6: <code>iload_1</code>		
-	(I)	(I,I)	7: <code>iload_0</code>		
-	(I,I)	(I,I)	8: <code>imul</code>		
-	(I)	(I,I)	9: <code>istore_1</code>		
-	()	(I,I)	10: <code>iinc 0, -1</code>		
-	()	(I,I)	11: <code>goto 2</code>		
-	()	(I,I)	14: <code>iload_1</code>		
-	(I)	(I,I)	15: <code>ireturn</code>	(I)	(I,I)

and we have completed the verification without error.

# Some of the Lattice of Types (Leroy, Figure 3)



```
class C {
}
class D extends C {
}
class E extends C {
}
```

## Merging Types

- The lattice represents an ordering relation on types
- The lattice is derived from the semantics of Java (and is based on the class hierarchy)
- Given any two types  $t_1$  and  $t_2$ , there is a least upper bound type,  $\text{lub}(t_1, t_2)$
- Given any type  $t$ , the length of the path from  $t$  to top,  $T$ , is finite (the well-foundedness property).

**The step in Sun's verification algorithm where types are merged is implemented as *lub*.**

**The finiteness property guarantees that Sun's algorithm will converge in a finite number of steps.**