



IBM Research

A Type-Based Approach for Data-Centric Synchronization

Frank Tip

joint work with Mandana Vaziri, Julian Dolby and Jan Vitek

Synchronization is difficult in OO programs

- **Possibility of data races or deadlock**
 - Low-level data race: 2 concurrent access to a location (at least one write) with no synchronization between them
- **Locking discipline: non-local reasoning**
- **Granularity of locking**
 - Even if every shared access is protected, data may still end up in an inconsistent state (e.g., high-level races)
- **Objective of this work**
 - Data-centric mechanism for synchronization constraints
 - Leverage OO structure
 - Automated synchronization inference
 - Lightweight ownership type system that guarantees serializability and enables separate compilation
 - improved support for linked data structures

Example: High-Level Data Race in java.util.Vector

```
public class Vector<E> {  
    protected Object[] elementData;  
    protected int elementCount;  
  
    ...  
  
    public Vector(Collection<? extends E> c) {  
        elementCount = c.size();  
        elementData = new Object[(int)Math.min  
((elementCount*110L)/100, Integer.MAX_VALUE)];  
        c.toArray(elementData);  
    }  
    ...  
}
```

Idea behind Data-Centric Synchronization

**Synchronization is about preserving
data consistency**

**So why not associate synchronization constraints
directly with data?**

Terminology

- synchronized block refers to set of locations accessed atomically
- part of a larger operation that preserves consistency

```
class BankAccount {  
    int checking, savings;  
    int transferCount;
```

atomic set S

unit of work on S
*preserves consistency
of S when executed
sequentially*


```
void transfer(int amount) {  
    synchronized(this) {  
        checking -= amount;  
        savings += amount;  
    }  
    transferCount++;  
    ...  
}
```

Language Construct: Atomic Sets

```
class Account {
    atomicset account;
    atomic(account) int checking;
    public void deposit(int n)
    { ... }
    public void withdraw(int n)
    { ... }
}
```

```
class Bank {
    atomicset logging;
    atomic(logging) Log log;
    atomic(logging) int logCount;

    void transfer(Account a,
                  Account b,
                  int n){
        log.add(a,b,n);
        a.withdraw(n);
        b.deposit(n);
        logCount++;
    }
}
```



unit of work for this.logging

Language Construct: `unitfor`

```
class Account {
    atomicset account;
    atomic(account) int checking;
    public void deposit(int n){ ... }
    public void withdraw(int n){ ... }
}

class Bank {
    atomicset logging;
    atomic(logging) Log log;
    atomic(logging) int logCount;

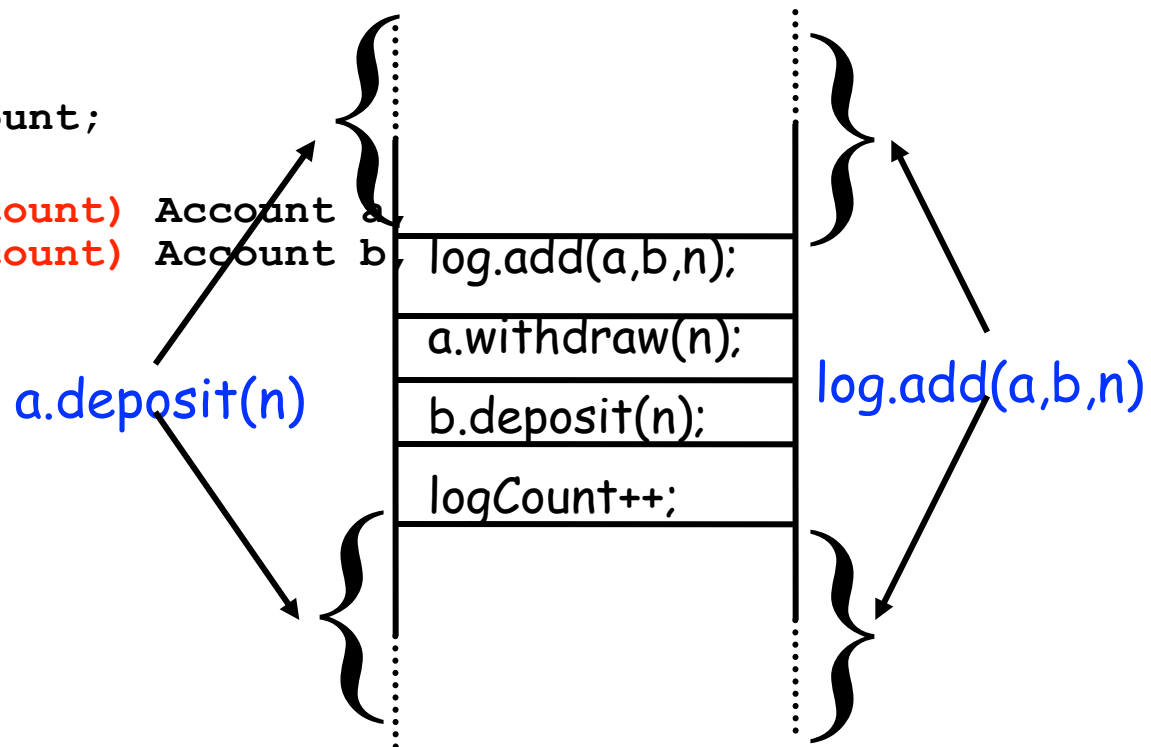
    void transfer(unitfor(account) Account a,
                  unitfor(account) Account b,
                  int n){
        log.add(a,b,n);
        a.withdraw(n);
        b.deposit(n);
        logCount++;
    }
}
```

Example: unitfor

```
class Account {
    atomicset account;
    atomic(account) int checking;
    public void deposit(int n){ ... }
    public void withdraw(int n){ ... }
}
```

```
class Bank {
    atomicset logging;
    atomic(logging) Log log;
    atomic(logging) int logCount;
```

```
void transfer(unitfor(account) Account a,
              unitfor(account) Account b,
              int n){
    log.add(a,b,n);
    a.withdraw(n);
    b.deposit(n);
    logCount++;
}
```



unit of work for a.account, b.account, this.logging

Preventing the High-Level Data Race in Vector

```
public class Vector<E> {  
    // atomicset L inherited from Collection  
    atomic(L) protected Object[] elementData;  
    atomic(L) protected int elementCount;  
  
    ...  
  
    public Vector(unifor(L) Collection<? extends E> c) {  
        elementCount = c.size();  
        elementData = new Object[(int)Math.min((elementCount*110L)/100,  
                                                Integer.MAX_VALUE)];  
        c.toArray(elementData);  
    }  
  
    ...  
}
```

Atomic sets that span multiple objects

- **Data structures often rely on other objects for their implementation**
 - The state in these objects belongs to that of the “owning” data structure
- **Examples**
 - LinkedList relies on LinkedList.Entry
 - HashMap relies on HashMap.Entry, a subtype of Map.Entry
 - Most collections rely on helper classes used to define iterators, views, ...

Aliasing and Internal Classes

- Mechanism for merging atomic sets that span multiple objects
- syntax: annotation `|X=this.Y|` on declarations of fields, local variables, allocation sites
- Meaning:
 - Atomic set `X` in the object pointed to by the annotated reference is merged with the atomic set `Y` in the object pointed to by `this`
 - All units of work for `Y` become units of work for `X` as well
- **internal classes**
 - must be aliased (i.e. atomic sets must be merged with owner)
 - access restricted to ensure that no reference to an internal class can leak outside of an atomic set
 - permits efficient implementation (no additional synchronization needed, because access is already protected by the class it is aliased with)

Example: Aliasing

the atomic set E in the object pointed to by header is merged with atomic set L in the object pointed to by this

```
class LinkedList extends ... {  
    atomic(L) private Entry header|E=this.L|;  
}
```

internal type: must be aliased, some restrictions on access, efficient implementation

```
internal class Entry {  
    atomicset E;  
    atomic(E) Object elem;  
    atomic(E) Entry next|E=this.E|;  
    atomic(E) Entry prev|E=this.E|;  
}
```

the atomic sets E in the objects pointed to by next/prev are merged with atomic set E in the object pointed to by this

Example: simplified version of java.util.LinkedList

- **abstract class AbsList**
 - declares `size()`, `iterator()`, `add(Object)`, `addAll(AbsList)`, `get(int)`
- **class LinkedList**
 - concrete implementation of AbsList
- **class Entry**
 - helper class for representing list-entries
- **interface ListIterator**
 - declares methods `hasNext()`, `next()`, `hasPrev()`, `prev()`, `set()`
- **class ListItr**
 - helper class that provides an implementation of ListIterator
- **class Client**
 - small multi-threaded client of LinkedList

Example, continued

```
class LinkedList extends AbsList {
    atomic(L) private Entry header|E=this.L|;

    public LinkedList() {
        header = new Entry|E=this.L|(null, null, null);
        header.next = header.prev = header;
    }

    ...

    public ListIterator iterator() {
        return new ListItr|I=this.L|(this, this.header, 0);
    }
    public boolean addAll(unitfor(L) AbsList c) {
        boolean modified = false;
        ListIterator e = c.iterator();
        while (e.hasNext()) {
            add(e.next());
            modified = true;
        }
        return modified;
    }
}
```

Example, continued

```
class ListItr implements ListIterator {
    atomicset I;
    atomic(I) private Entry lastReturned|E=this.I|;
    atomic(I) private Entry next|E=this.I|;
    atomic(I) private int nextIndex;
    atomic(I) final LinkedList list|L=this.I|;
    atomic(I) final Entry header|E=this.I|;

    ListItr(LinkedList l|L=this.I|, Entry h|E=this.I|, int index) {
        list = l;
        header = h;
        lastReturned = header;
        if (index < 0 || index > list.size())
            throw new IndexOutOfBoundsException();
        next = header.next;
        for (nextIndex = 0; nextIndex < index; nextIndex++)
            next = next.next;
    }

    // hasNext(), next(), hasPrev(), prev()

    public void set(Object o) {
        if (lastReturned == header)
            throw new IllegalStateException();
        lastReturned.elem = o;
    }
}
```

Example, continued

```
public class Client {
    public static void main(String[] args) throws Throwable {
        final AbsList x = new LinkedList();
        final AbsList y = new LinkedList(); y.add("a"); y.add("a");
        final AbsList z = new LinkedList(); z.add("b"); z.add("b");
        Thread t1 = new Thread(){
            public void run(){ x.addAll(y); }
        };
        Thread t2 = new Thread(){
            public void run(){ x.addAll(z); }
        };
        t1.start(); t2.start();
        t1.join(); t2.join();
        ListIterator it;
        for (it = x.iterator(); it.hasNext();){
            Object o = it.next();
            if (o.equals("b")) it.set("c");
        }
        // can print aacc or ccaa, but not acac, caca, caac, acca
        for ( ; it.hasPrev(); )
            System.err.println(it.prev());
    }
}
```


Type System (Details in Technical Report)

- **Types are of the form: `C` or `C|a=this.b|`**
 - object of type `C` is in charge of its own synchronization
 - object of type `C|a=this.b|` has its synchronization managed by atomic set `this.b`
- **internal types**
 - must have alias information
 - internal object with alias `|a=this.b|` must be accessed from within a unit of work for `b`
 - (can be implemented efficiently, because synchronization is already ensured by owner)
- **Type rules ensure consistency of alias annotations and encapsulation**
- **Key properties establish atomic-set serializability:**
 - Units of work on each atomic set must be serializable.
 - Since units of work on `S` are mutually exclusive, and no access to `S` happens outside of a unit of work on `S`, all units of work on `S` happen serially.

Formalization

- **Give dynamic semantics as small-step operational semantics, including heaps and threads**
- **Define configuration as a heap and a series of threads**
 - each thread has a thread ID and a stack
- **Define well-formedness conditions on configurations**
 - configuration is well-formed if its heap and threads are well-formed
 - heap is well-formed if, for every object, the values of its fields are subtypes of the declared types of those fields
 - thread is well-formed if all the frames in its stack are well-typed
- **Prove type soundness (progress and preservation)**
- **Prove that internal objects with alias $|a=this.b|$ are only accessed from within a unit of work for b**
- **Prove atomic-set-serializability**

Implementation

- **source-to-source translator implemented using Eclipse refactoring infrastructure**
 - atomic set annotations entered as Java comments
 - implementation handles Java subset (no generics, inner classes)
 - incl. annotations for arrays (e.g. `|this.M[]F=this.M|`)
 - two translation options supported:
 1. based on `java.util.concurrent.locks.ReentrantLock`
 2. based on lock ordering & standard synchronized mechanism
- **translation involves**
 - create lock field `$lock_S` in any class C such that atomic set S is present
 - for each class with a lock field `$lock_S`, generate methods `takeLockForS()`, `tryLockForS()`, and `releaseLockForS()`
 - transform object allocations to set locks
 - transform units of work to acquire all needed locks
- **doing a build invokes the translator on the current project**
 - translated sources saved in a separate project in the workspace

AbsList Translated

```
abstract class AbsList implements atomicsets.Atomic {
    protected Lock $lock_L;

    public final void takeLockForL() {
        if ($lock_L != null)
            $lock_L.lock();
    }
    public final boolean tryLockForL() {
        return $lock_L == null || $lock_L.tryLock();
    }
    public final void releaseLockForL() {
        if ($lock_L != null)
            $lock_L.unlock();
    }
    public final Lock getLockForL() {
        return $lock_L;
    }
    public AbsList setLockForL(Lock l) {
        $lock_L = l; return this;
    }
    /*atomic(L)*/ int size = 0;

    public int size() {
        try {
            this.takeLockForL();
            { return size; }
        } finally {
            this.releaseLockForL();
        }
    }
}
```

generated lock field for atomic set L

*generated methods for
manipulating the lock
associated with atomic set L*

translated size() method

LinkedList.iterator() Translated

```
class LinkedList extends AbsList {  
    ...  
    public ListIterator iterator() {  
        try {  
            this.takeLockForL();  
            {  
                return (ListIterator)  
                    new ListItr(this, this.header, 0).setLockForI(this.getLockForL());  
            }  
        } finally {  
            this.releaseLockForL();  
        }  
    }  
}
```



translated allocation (aliased)

LinkedList.addAll() Translated

```
public boolean addAll(/*unitfor(L)*/ AbsList c) {
    boolean $repeat = true;
    do {
        try {
            this.takeLockForL();
            if (c.tryLockForL())
                try {
                    $repeat = false;
                    {
                        boolean modified = false;
                        ListIterator e = c.iterator();
                        while (e.hasNext()) {
                            add(e.next());
                            modified = true;
                        }
                        return modified;
                    }
                } finally {
                    c.releaseLockForL();
                }
            } finally {
                this.releaseLockForL();
            }
        } while ($repeat);
        throw new Error();
    }
}
```

atomically obtain locks for multiple atomic sets

Client Translated

```

public class Client {
    public static void main(String[] args) throws Throwable {
        final AbsList x = new LinkedList().setLockForL(new ReentrantLock());
        final AbsList y = new LinkedList().setLockForL(new ReentrantLock());
        y.add("a"); y.add("a");
        final AbsList z = new LinkedList().setLockForL(new ReentrantLock());
        z.add("b"); z.add("b");
        Thread t1 = new Thread(){
            public void run(){ x.addAll(y); }
        };
        Thread t2 = new Thread(){
            public void run(){ x.addAll(z); }
        };
        t1.start(); t2.start();
        t1.join(); t2.join();
        ListIterator it;
        for (it = x.iterator(); it.hasNext();){
            Object o = it.next();
            if (o.equals("b")) it.set("c");
        }
        // can print aacc or ccaa, but not acac, caca, caac, acca
        for ( ; it.hasPrev(); )
            System.err.println(it.prev());
    }
}

```

object allocation, not aliased

Experiment #1: Java Collections Framework

- **selected several classes from Java Collections Framework**
 - ArrayList, LinkedList, HashMap, HashSet, TreeMap, LinkedHashMap, LinkedHashSet
 - along with any types in java.util on which they depend
 - 63 types, 10,860 LOC
- **manually refactored these classes to eliminate generics and nested classes**
 - not yet supported by our implementation
- **then introduced atomic sets**
 - one atomic set for each of 5 subhierarchies, which includes all instance fields
 - uniform annotations on “bulk” methods such as addAll()
 - alias annotations to relate entries, iterators, views, etc. to their “owner”
 - only one class could be made internal (LinkedList.Entry)

Experiment #1: Java Collections Framework

- **selected several classes from Java Collections Framework**
 - ArrayList, LinkedList, HashMap, HashSet, TreeMap, LinkedHashMap, LinkedHashSet
 - along with any types in java.util on which they depend
 - 63 types, 10,860 LOC
- **manually refactored these classes to eliminate generics and nested classes**
 - not yet supported by our implementation
- **then introduced atomic sets**
 - one atomic set for each of 5 subhierarchies, which includes all instance fields
 - uniform annotations on “bulk” methods such as addAll()
 - alias annotations to relate entries, iterators, views, etc. to their “owner”
 - only one class could be made internal (LinkedList.Entry)

Annotation Overhead for classes from JCF

annotation type	number of annotations
atomicset	0
atomic class	5
atomic	0
unitfor	55
alias	330
array object	24
array element	16
TOTAL	430

- ~1 annotation per 25 lines of code

Experiment #2: SPECjbb2005

- **widely used multi-threaded performance benchmark**
- **simulates companies, warehouses, customers, orders, etc.; performs a series of 4-minute runs with increasing number of warehouses**
- **has a built-in performance mechanism for measuring throughput (transactions/second)**
- **uses both `synchronized` and `wait()` / `notify()`**
- **7,891 LOC**

Refactoring SPECjbb2005: Some Observations

- **We analyzed the shared state & synchronization in SPECjbb, and introduced atomic sets in a way that ensures that all access to shared state is properly synchronized.**
- **synchronization seems inconsistent in several places**
 - some shared fields, e.g., `Customer.creditLimit` have synchronized accessors, but others, e.g., `Customer.address` have unsynchronized accessors
 - some methods, e.g. `TreeMapDataStorage.deleteFirstEntities()` should logically be executed atomically, but there is no synchronization to enforce this
- **redundant synchronization**
 - e.g., of fields that are written only once during execution of the constructor and that could be made final
- **`wait()` and `notify()` used to implement barriers that coordinate the threads of the multiple warehouses.**
 - `wait()/notify()` must be used with care when atomic sets are introduced.
- **ownership issues: some classes in SPECjbb rely on collections to store data**
 - e.g., `TreeMapDataStorage` relies on a `TreeMap` to store its data
 - achieved by aliasing their atomic sets

Annotation Overhead for SPECjbb2005

annotation type	number of annotations
atomicset	1
atomic class	14
atomic	25
unitfor	0
alias	8
array object	0
array element	1
TOTAL	49

- these annotations replace 125 occurrences of `synchronized` in the original code
- 25 occurrences of `synchronized` remain that are related to the use of `wait()/notify()`
- ~1 annotation per 160 lines of code

SPECjbb2005: Performance Experiments

	1	2	3	4	5	6	7	8
original	9654	13341	13566	12854	12914	12875	13951	14063
naive	5588	8068	8074	8030	8106	8205	8159	8248
tuned	6295	8394	8498	8767	8945	9380	9277	9403

- **tuned version: refactored code to make fields final, so our compiler can avoid inserting synchronization**
- **bottom line:**
 - naive implementation achieves approx. 60% of performance of original
 - tuned implementation achieves approx. 67% of performance of original

Related Work (not a complete list)

- **previous work on atomic sets**
 - language proposal, correctness criterion [Vaziri et al. POPL'06]
 - dynamic detection of a-set serializability violations [Hammer et al. ICSE'08]
 - static detection of a-set serializability violations [Kidd et al. VMCAI'09]
- **data groups: abstract representation of groups of fields for modular reasoning**
 - Leino et al. [OOPSLA'98, PLDI'02]
- **type systems for ensuring race-freedom and atomicity**
 - Flanagan, Freund et al. [ESOP'99, PLDI'00, PLDI'03, TOPLAS'08, ...]
- **ownership type systems**
 - Noble, Vitek et al. [ECOOP'98, OOPSLA'98, OOPSLA'99], Banerjee, Naumann [POPL'02], Clarke, Drossopoulou [OOPSLA'02], ...
 - unlike traditional owner-as-dominator type systems, there is no single access point
- **lock inference for atomic sections**
 - Cherem et al. [PLDI'08], McCloskey et al. [POPL'06]

Conclusions

- **type-based data-centric approach to synchronization**
(our original approach required whole-program analysis)
 - correctness property: a type-correct program is atomic-set serializable
 - enables separate compilation
 - handles linked data structures naturally; expressive enough for real code
- **annotation overhead similar to that of previous type systems that guarantee race-freedom or atomicity**
 - but those previous type systems were in addition to existing synchronization
- **preliminary performance experiments indicate reasonable performance**
 - tuned version SPECjbb achieves 67% of throughput of the original code
 - many opportunities for optimization

Future Work

- **improve implementation**
 - improve performance of generated code (e.g., use ReadWriteLocks)
 - automatically infer annotations of most internal types
 - investigate lock-free implementation
- **design analysis to detect possible deadlock**
- **extend type system with effects on methods for better code generation**
 - e.g., “reads this.L”, “writes this.L”

RACKI ID

Example: Low-Level Data Race

```
public class Example extends Thread {  
    private static int cnt = 0; // shared state  
    public void run() {  
        int y = cnt;  
        cnt = y + 1;  
    }  
  
    public static void main(String[] args) {  
        Thread t1 = new Example();  
        Thread t2 = new Example();  
        t1.start();  
        t2.start();  
    }  
}
```

Example: Low-Level Data Race

```
private static int cnt = 1; // shared state
```

thread #1

```
public void run() {  
    int v = cnt;    y=0  
    cnt = y + 1;  
}
```

thread #2

```
public void run() {  
    int v = cnt;    y=0  
    cnt = y + 1;  
}
```

Example: High-Level Data Race in java.util.Vector

```

public class Vector<E> {
    protected Object[] elementData;  ———→ 3
    protected int elementCount;      ———→ [null,null,null]

    public Vector(Collection<? extends E> c) {
        elementCount = c.size();
        elementData = new Object[(int)Math.min((elementCount*110L)/100,
                                                Integer.MAX_VALUE)];
        c.toArray(elementData);
    }
}

c ———→ []

```

thread #1

```

elemCount = c.size();
elemData = new Object[ ...elemCount...];
c.toArray(elemData);

```

thread #2

```
c.clear();
```

Type System

- **Subtyping:**

$$\frac{C <: D}{C|a=this.b| <: D|a=this.b|}$$
- **Rules for object creation**
 - An internal object must be created with alias information
 - ‘C has a’ means class C defines or inherits atomic set a

(T-New-Raw)

$E(x)=C, \quad C \text{ not internal}$

$E \text{ d } x = \text{new } C()$

(T-New-Aset)

$E(x)= C|a=this.b|, C \text{ has } a, \quad E(\text{this}) \text{ has } b$

$E \text{ d } x = \text{new } C|a=this.b|()$

Type System (2)

Rules for casts

- Casts are explicit in the formal language
- Alias information of non-internal objects can be erased

(T-Cast-Plain)

$$E(x)=D, \quad E(y)=C, \quad D <: C$$

$$E \text{ d } y = (C) x$$

(T-Cast-Aset)

$$E(x)=D|a=\text{this}.b|, \quad E(y)=C|a=\text{this}.b|$$

$$C \text{ has } a, \quad E(\text{this}) \text{ has } b, \quad D <: C$$

$$E \text{ d } y = (C|a=\text{this}.b|) x$$

(T-Cast-Off)

$$E(x)=C|a=\text{this}.b|, \quad C \text{ not internal}, \quad E(y)=C$$

$$E \text{ d } y = (C) x$$

Type System (3)

■ Rule for method call

- **adapt** is a predicate used to unify alias annotations, to “adapt” an alias annotation to the context of the receiver
- When adapt is undefined, type rules fail, preventing internal objects from being used in a context in which their alias annotation is not well-defined.

$$\text{adapt}(C, t) = C$$

$$\text{adapt}(C|a=\text{this}.b|, D|b=\text{this}.c|) = C|a=\text{this}.c|$$

(T-Call)

$$\frac{E(y)=t_y, \text{typeof}(t_y.m)=\bar{t} \rightarrow t, E(\bar{z})=\bar{t}_z \quad \bar{t}_z <: \text{adapt}(\bar{t}, t_y), t' = \text{adapt}(t, t_y), E(x)=t'}{E \text{ d } x = y.m(\bar{z})}$$

Type System (4)

- Rules for field selection & update

(T-Select)

$$\begin{array}{l} E(x)=t_x, E(y)=t_y, \text{typeof}(t_y.f)=t \\ t_x <: \text{adapt}(t, t_y), \quad (t_y.f \text{ is atomic} \Rightarrow y = \text{this}) \end{array}$$

$$E \text{ d } x = y.f$$

(T-Update)

$$\begin{array}{l} E(x)=t_x, E(y)=t_y, \text{typeof}(t_x.f)=t \\ t_y <: \text{adapt}(t, t_x), \quad (t_x.f \text{ is atomic} \Rightarrow x = \text{this}) \end{array}$$

$$E \text{ d } x.f = y$$

Example: Aliasing

```

abstract class AbsList {
    atomicset L;
    atomic(L) int size = 0;
    public int size(){ return size; }
    public abstract ListIterator iterator();
    public abstract void add(Object o);
    public abstract boolean addAll(unitfor(L) AbsList c);
    public abstract Object get(int i);
}

class LinkedList extends AbsList {
    atomic(L) private Entry header|E=this.L|;
    public LinkedList(){
        header = new Entry|E=this.L|(null, null, null);
        header.next = header.prev = header;
    }
    public void add(Object o){
        Entry newEntry|E=this.L| = new Entry|E=this.L|(o, header, header.prev);
        newEntry.prev.next = newEntry;
        newEntry.next.prev = newEntry;
        size++;
    }
}

internal class Entry {
    atomicset E;
    atomic(E) Object elem;
    atomic(E) Entry next|E=this.E|;
    atomic(E) Entry prev|E=this.E|;
}

```

the atomic set E in the object pointed to by header is merged with atomic set L in the object pointed to by this

internal type: must be aliased, some restrictions, efficient implementation

the atomic sets E in the objects pointed to by next/prev are merged with atomic set E in the object pointed to by this