

# Concurrent Programming in Java




Thursday, August 27, 2009



ECS 2009

## Concurrency Models I

- Heavyweight tasks execute in their own address space
- Lightweight tasks run in the same address space
- A task is disjoint if it does not communicate with or affect the execution of any other task



Thursday, August 27, 2009

## Concurrency Models II

- Java supports threads
  - ▶ Threads execute within a single JVM
  - ▶ **Native threads** map a single Java thread to an OS thread
  - ▶ **Green threads** adopt the thread library approach
  - ▶ **M-on-N** threads are a mixture of the above (M green threads scheduled on N native threads)
  - ▶ On a multiprocessor, native threads are needed to get true parallelism

Thursday, August 27, 2009


## Concurrency Models III

- There are various ways in which concurrency can be introduced
  - ▶ API for explicit thread creation or thread forking
  - ▶ high-level language construct, e.g. PAR (occam), tasks (Ada), or processes (Modula)
- Integration with object-oriented programming, various models:
  - ▶ **asynchronous method calls**
  - ▶ **early return** from methods
  - ▶ **futures**
  - ▶ **active objects**
- Java adopts the active object approach

Thursday, August 27, 2009

## Concurrency Models IV

- Communication and Synchronization
  - approaches broadly classified as shared-variable or message passing
  - many different models, a popular one is a monitor
  - a monitor can be considered as an object where each of its operation executes in **mutual exclusion**



Thursday, August 27, 2009

## Concurrency Models V

- **Condition Synchronization**
  - expresses a constraint on the ordering of execution of operations, e.g., data cannot be taken from a buffer until data has been put in it
- Monitors provide condition variables with three operations which can be called when the lock is held
  - **wait**: an unconditional suspension of the calling thread (the thread is placed on a queue associated with the condition variable)
  - **notify**: one thread is taken from the queue and re-scheduled for execution (it must reclaim the lock first)
  - **notifyAll**: all suspended threads are re-scheduled
  - **notify** and **notifyAll** have no effect if no threads are suspended on the condition variable

Thursday, August 27, 2009

# Kinds of Synchronization

- **Cooperative**

- ▶ Task A must wait for task B to complete some activity before A can continue executing, e.g., producer-consumer
- ▶ **wait/notify**

- **Competitive**

- ▶ Two or more tasks must use a resource that cannot be simultaneously accessed, e.g., a shared counter
- ▶ **synchronized**

Thursday, August 27, 2009

# Liveness and Deadlock

- **Liveness** is a characteristic that a program may or may not have

- ▶ In sequential code, it means the program will eventually complete
- ▶ In a concurrent environment, a task can easily lose its liveness
- ▶ If all tasks lose their liveness, it is called **deadlock**

Thursday, August 27, 2009

## Concurrency in the real world...

- Concurrency is also a source of problems
  - Windows 2000: Concurrency errors most common defects among detectable errors
  - Windows 2000: Incorrect synchronization and protocol errors most common coding errors
  - Windows 2003: Synchronization errors second only to buffer overruns
  - Race conditions create security vulnerabilities
  - Concurrent programs are hard to test because they are non-deterministic
    - *Bugs are hard to reproduce because there are exponentially many possible interleavings that often only manifest on deployed*



Thursday, August 27, 2009

## Data Races

- A *race condition* occurs if two threads access a shared variable at the same time and at least one of the accesses is a write
- Consider the following program when multiple threads are call **deposit()** in parallel:

```
class Account {
    private int bal;

    void deposit(int n) {
        int j = bal;
        bal = j + n;
    }
}
```

Thursday, August 27, 2009

## Lock-based Synchronization

- Monitors must be used to protect every shared location access

```
a = x.a; b = x.b;  ||  x.a = 1; x.b = 2;
```

- Locks must be held before every read/write to a shared location

```
synchronized(x) {a = x.a; b = x.b;}  
||  
synchronized(x) {x.a
```

- When multiple locks are used, a lock acquisition protocol must be adhered to by the application to avoid deadlocks

```
synchronized(x){ synchronized(y) {...}}  
||  
synchronized(y){ synchronized(x) {...}}
```

Thursday, August 27, 2009

## Puzzlers #1

- How many different values can there be for local variables a and b?

```
a = x.a; b = x.b;  ||  x.a = 1; x.b = 2;
```

Thursday, August 27, 2009

## Puzzlers #2

- Do the following programs have data races? Are they non-deterministic?

```
long a = x.a;      ||    long b = x.a;
```

```
x.i=1;int i=x.i;   ||    x.i=1;int i=x.i;
```

```
x.a = MAX_LONG
```

```
x.a = 0           ||    long a = x.a;  x.a = (a==MAX_LONG)? 0 : a;
```

Thursday, August 27, 2009

## Puzzlers #3

- Does the following program have data race?

```
interface INC { void inc(); }
```

```
class Int implements INC { int i; void inc() { i++; } }
```

```
class SyncInt implements INC {
    INC val;
    SyncInt(INC v){ val = v; }
    synchronized inc() { val.inc(); }
}
```

```
...
```

```
INC i = new Int();
```

```
INC[] arr=new INC[]{new SyncInt(i),new SyncInt(i),new SyncInt(i)};
```

```
arr[0].inc();      ||    arr[1].inc();      ||    arr[2].inc();
```

Thursday, August 27, 2009

## Puzzlers #3

- Does the following program have concurrency problem?

```
class LL {
    int i;
    LL next;
    LL(LL n,int v){next=n;i=v;}
    synchronized swap() {
        synchronized(next) { int t=next.i; next.i= i; i=t; }
    }
}
```

```
LL a=new LL(null,0), b=new LL(a,1); a.next=b;
```

```
a.swap();    ||    b.swap();
```

Thursday, August 27, 2009

## Puzzlers #3

- A fix?

```
class LL {
    static int K;
    private final int id = K++;
    int i;
    LL next;
    LL(LL n,int v){next=n;i=v;}
    private void _swap() {int t=next.i; next.i= i; i=t;}
    void swap() {
        if (this.id > next.id)
            synchronized (this) {synchronized(next) { _swap(); }
        }
        else
            synchronized (next) {synchronized(this) { _swap(); }
        }
    }
}
```

Thursday, August 27, 2009



## Puzzlers #3

- A fix?

```
class LL {
    static int K;
    private final int id;
    int i; LL next;

    LL(LL n, int v) {
        next = n; i = v;
        synchronized (LL.class) { id = K++; }
    }

    private void _swap() { int t = next.i; next.i = i; i = t; }
    void swap() {
        if (this.id > next.id)
            synchronized (this) { synchronized (next) { _swap(); } }
        else
            synchronized (next) { synchronized (this) { _swap(); } }
    }
}
```

Thursday, August 27, 2009

## Puzzlers #4

- The following idiom occurs frequently in library code. Why is it?

```
final class SyncTree {

    Node left, right;

    private final Object lock = new Object();

    public balance() {

        synchronized (lock) {

            ....
        }
    }
}
```

Thursday, August 27, 2009

## Puzzlers #5

- Is this a correct solution to synchronization problems?

```
class Big {
    static final Object Lock = new Object();
}

class LL {
    int i;
    LL next;
    LL(LL n,int v){next=n;i=v;}
    void swap() {
        synchronized(Big.Lock) { int t=next.i; next.i= i; i=t; }
    }
}

LL a=new LL(null,0), b=new LL(a,1); a.next=b;

a.swap();    ||    b.swap();
```

Thursday, August 27, 2009

## Puzzlers #6

- Does the following program have a data race? A concurrency problem?

```
class Big { static final Object Lock = new Object(); }

class LL {
    int i; LL next; LL(LL n,int v){next=n;i=v;}

    void swap() {
        int t = 0;
        synchronized(Big.Lock) { t=next.i; }
        synchronized(Big.Lock) { next.i= i; }
        synchronized(Big.Lock) { i=t; }
    }
}
```

Thursday, August 27, 2009

## Lessons

- To reason about concurrency one must understand all interleaving of operations performed by each thread
- Not all high-level commands are atomic
- Aliasing makes it hard to determine which values are shared, thus one may fail to acquire the right lock (or locks)
- Lock acquisition protocols must be followed to avoid deadlocks
- Library classes must protect themselves from clients by implementing their own synchronization
- Oversynchronization is always safe but decreases concurrency (possibly to the point of making it meaningless?)
- Definition of data race too low level to catch all errors

Thursday, August 27, 2009

## Concurrency in Java

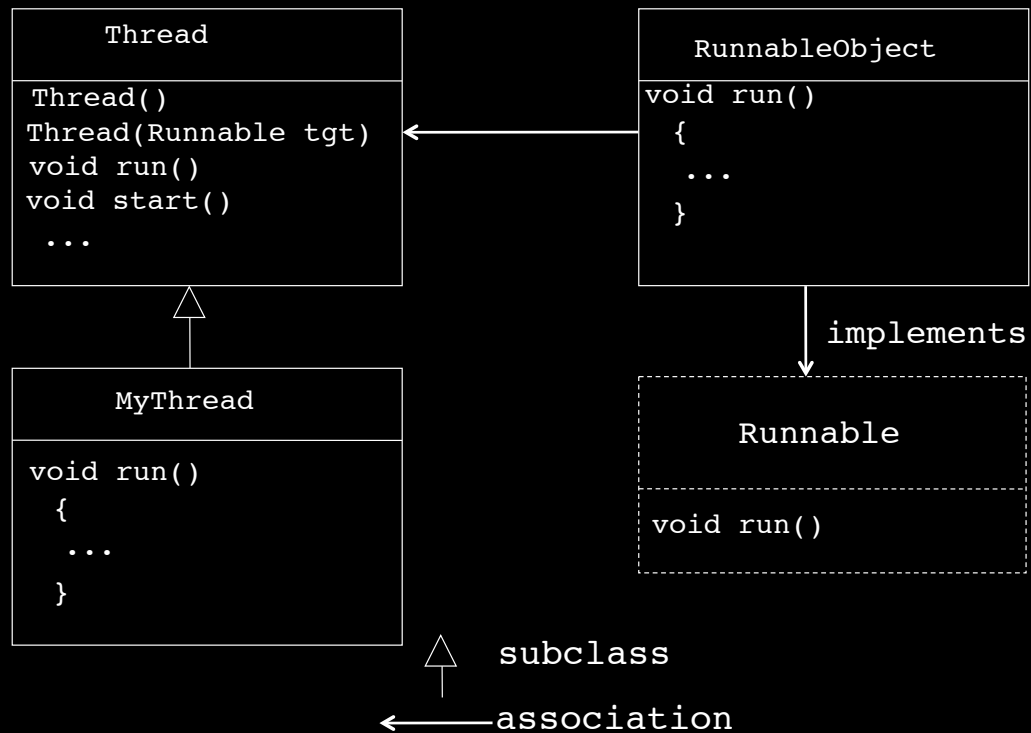
- Java has a predefined class `java.lang.Thread` which provides the mechanism by which threads are created
- However to avoid all threads extending `Thread`, it also has a standard interface

```
public interface Runnable {  
    public void run();  
}
```

- Hence, a class which wishes to express concurrent execution implements this interface
- Threads do not begin their execution until `start` is called

Thursday, August 27, 2009

# Threads in Java



Thursday, August 27, 2009

# Communication in Java

- Via reading and writing to data encapsulated in shared objects protected by monitors
- Every object is implicitly derived from the `Object` class which defines a mutual exclusion lock
- Methods in a class can be labeled as **synchronized**, this means that they can only be executed if the lock can be acquired (this happens automatically)
- The lock can also be acquired via a **synchronized statement** which names the object
- A thread can **wait** and **notify** on a single anonymous condition variable

Thursday, August 27, 2009

# The Thread Class

```
public class Thread implements Runnable {  
    public Thread();  
    public Thread(String name);  
    public Thread(Runnable target);  
    public Thread(Runnable target, String name);  
    public Thread(Runnable target, String name,  
                  long stackSize);  
  
    public void run();  
    public void start();  
    ...  
}
```

Thursday, August 27, 2009

## Thread Identification

- The identity of the currently running thread can be found using the `currentThread` method
- This has a static modifier, which means the method can always be called using the `Thread` class

```
public class Thread implements Runnable {  
    ...  
    public static Thread currentThread();  
}
```

Thursday, August 27, 2009

## A Thread Terminates:

- when it completes execution of its **run** method either normally or as the result of an unhandled exception
- via a call to its **stop** method — the **run** method is stopped and the thread class cleans up before terminating the thread (releases locks and executes any finally clauses)
  - the thread object is now eligible for garbage collection.
  - **stop** is inherently unsafe as it releases locks on objects and can leave data in inconsistent states; (deprecated; should not be used)
- via a call to its **destroy** method — **destroy** terminates the thread without any cleanup (deprecated)

Thursday, August 27, 2009

## Daemon Threads

- Threads can be of two types: **user** threads or **daemon** threads
- Daemon threads provide general services and never terminate
- When all user threads have terminated, daemon threads will be terminated by the virtual machine on shutdown
- The **setDaemon** method must be called before calling **start**

```
public class Thread implements Runnable {  
    public void destroy();           // DEPRECATED  
    public final boolean isDaemon();  
    public final void setDaemon();  
    public final void stop();        // DEPRECATED  
}
```

Thursday, August 27, 2009

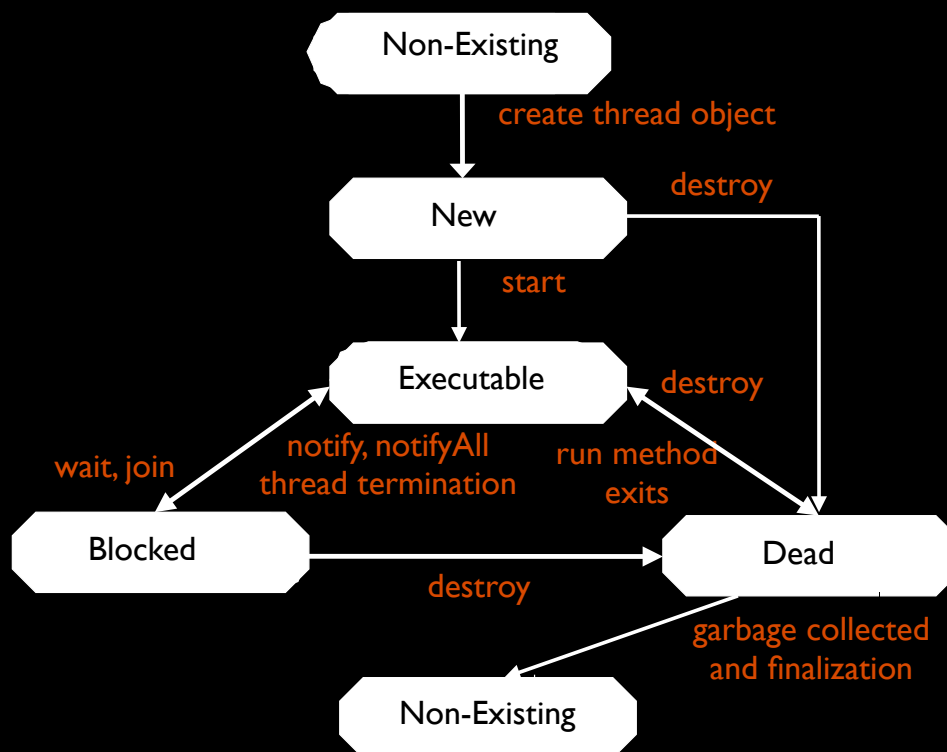
## Joining

- One thread can wait (with or without a timeout) for another thread (the target) to terminate by issuing the **join** method call on the target's thread object
- The **isAlive** method allows a thread to determine if the target thread has terminated

```
public class Thread implements Runnable {
    public final native boolean isAlive();
    public final void join() throws InterruptedException;
    public final void join(long ms) throws InterruptedException;
    public final void join(long millis, int nanos) throws In
```

Thursday, August 27, 2009

## Java Thread States



Thursday, August 27, 2009

## Summary

- A thread is created when a **Thread** object is created
- At this point, the thread is not executable, it is in the **new** state
- Once **start** has been called, the thread is eligible for execution
- If the thread calls **wait** on an Object, or calls **join** on another thread object, the thread becomes **blocked** and no longer eligible
- It becomes executable if an associated **notify** is called by another thread, or if the target thread of the **join** is **dead**
- it enters the **dead** state, if the **run** method exits (normally or unhandled exception) or because **destroy** has been called. In the latter case, the thread will not execute any **finally** clauses; it may leave other objects locked

Thursday, August 27, 2009

## Synchronized Methods

- A mutual exclusion lock is associated with each object. It can't be accessed directly by the application but is affected by
  - ▶ the method modifier **synchronized**
  - ▶ block synchronization
- When a method is labeled as **synchronized**, the method can only execute once the system has the lock
- Hence, synchronized methods have mutually exclusive access to the data encapsulated by the object, if that data is only accessed by other synchronized methods
- Non-synchronized methods do not require the lock and, therefore, can be called at any time

Thursday, August 27, 2009



# Example of Synchronized Methods

```
class SharedInteger {  
    private int data;  
    SharedInteger(int val) { data = val; }  
    synchronized int read() { return data;}  
    synchronized void write(int val) { data = val; }  
    synchronized void incrementBy(int by) { data += by; }  
}  
  
SharedInteger shi = new SharedInteger(42);
```

Thursday, August 27, 2009

## Block Synchronization

- A mechanism where a block can be labeled as synchronized
- The **synchronized** keyword takes as a parameter an object whose lock the system needs to obtain before it can continue
- Synchronized methods are effectively implementable as

```
public int read() {  
    synchronized(this) {  
        return theData;  
    }  
}
```

- **this** is the Java mechanism for obtaining the current object

Thursday, August 27, 2009

## Warning

- In its full generality, **synchronized blocks** can undermine one of the advantages of monitor-like mechanisms, that of encapsulating synchronization constraints associate with an object into a single place in the program
- This is because it isn't possible to understand the synchronization associated with an object by just looking at the object itself when other objects can name that object in a synchronized statement
- However with careful use, this facility allows more expressive synchronization constraints to be programmed

Thursday, August 27, 2009

## Accessing Synchronized Data

- Consider a simple class which implement a two-dimensional coordinate that is to be shared between two or more threads
- This class encapsulates two integers
- Writing is simple, the **write** method can be synchronized
- The constructor can be assumed not to have any synchronization

```
class Coord {
    Coord(int x,int y) { x_=x; y_=y; }
    synchronized void write(int x,int y) {x_=x; y_=y;}
    private int x_, y_;
}
```

Thursday, August 27, 2009

## Shared Coordinate

- How to read the value of the coordinates?
- Methods return a single value, parameters are passed by value
- Consequently, it is not possible to have a single read method which returns both the **x** and the **y** values
- If two synchronized functions are used, **readX** and **readY**, the value of the coordinate can be written between calls. The result will be an inconsistent value of the coordinate

Thursday, August 27, 2009

## Solution I

- Return a new coordinate, which can be accessed freely

```
class Coord {  
    synchronized Coord read() { return new Coord(x, y); }  
    int readX() {return x;}    int readY() { return y; }  
}
```

- The result is only a snapshot of the shared `Coord`, which might be changed by another thread right after the **read** has returned
- The individual field values will be consistent
- Once the coordinate has been used, it can be discarded and made available for garbage collection
- If efficiency is a concern, avoiding unnecessary object creation and garbage collection is appropriate

Thursday, August 27, 2009

## Solution 2

- Assume the client thread will use synchronized blocks to obtain atomicity

```
class Coord {
    ...
    synchronized void write(int x, int y) { ... }

    int readX() { return x; } // not synchronized
    int readY() { return y; } // not synchronized
}

Coord p = new Coord(0,0);

synchronized(p) {
    Coord p2 = new Coord(p.readX(), p.readY());
}
```

Thursday, August 27, 2009

## Waiting and Notifying

- To obtain conditional synchronization requires the methods provided in the predefined object class
- These methods require the current thread to hold the object lock
- If called without the lock, the unchecked exception `IllegalMonitorStateException` is thrown
- The `wait` method always blocks the calling thread and releases the lock associated with the object

```
class Object {
    final void notify();
    final void notifyAll();

    final void wait() throws InterruptedException;
    final void wait(long millis) throws InterruptedException;
    final void wait(long millis, int nanos) throws InterruptedE
```

Thursday, August 27, 2009

## Important Notes

- The **notify** method wakes up one waiting thread; the one woken is not defined by the Java language
- **notify** does not release the lock; hence the woken thread must wait until it can obtain the lock before proceeding
- To wake up **all** waiting threads requires use of **notifyAll**
- If no thread is waiting, then **notify/notifyAll** are no-ops

Thursday, August 27, 2009

## Thread Interruption

- A waiting thread can also be awoken if it is interrupted by another thread
- In this case the **InterruptedException** is thrown

Thursday, August 27, 2009

## Condition Variables I

- There are **no** explicit condition variables in Java
- When a thread is awoken, it cannot assume that its condition is true, as all threads are potentially awoken irrespective of what conditions they were waiting on
- For some algorithms this limitation is not a problem, as the conditions are mutually exclusive,
- E.g., the bounded buffer traditionally has two condition variables: **BufferNotFull** and **BufferNotEmpty**
- If a thread is waiting for one condition, no other thread can be waiting for the other condition
- One would expect that the thread can assume that when it wakes, the buffer is in the appropriate state

Thursday, August 27, 2009

## Condition Variables II

- This is not always the case; Java makes no guarantee that a thread woken from a wait will gain immediate access to lock
- Another thread could call the put method, find that the buffer has space and inserted data into the buffer
- When the woken thread eventually gains access to the lock, the buffer will again be full
- Hence, it is usual for threads to re-evaluate their guards

Thursday, August 27, 2009

## Bounded Buffer

```
class BoundedBuffer {
    private final int buffer[];
    private int first, last, numberInBuffer;
    private final int size;
    BoundedBuffer(int len){ buffer = new int[size = len]; }

    synchronized void put(int i) throws InterruptedException{
        while (numberInBuffer == size) wait();
        numberInBuffer++;
        buffer[last = (last+1)%size] = i;
        notifyAll();
    }

    synchronized int get() throws InterruptedException {
        while (numberInBuffer == 0) wait();
        numberInBuffer--;
        notifyAll();
        return buffer[first = (first+1)%size];
    }
}
```

Thursday, August 27, 2009

## Class Exercise

- How would you implement a semaphore using Java?

Thursday, August 27, 2009

# Summary I

- Errors in communication and synchronization cause working programs to suddenly suffer from deadlock or livelock
- The Java model revolves around controlled access to shared data using a monitor-like facility
- The monitor is represented as an object with synchronized methods and statements providing mutual exclusion
- Condition synchronization is given by the wait and notify method
- True monitor condition variables are not directly supported by the language and have to be programmed explicitly