

Java 1.5 Concurrency Utilities

- Comprehensive support for general-purpose concurrent programming; partitioned into three packages:
 - ▶ `java.util.concurrent` — support common concurrent programming paradigms, e.g., various queuing policies such as bounded buffers, sets and maps, thread pools
 - ▶ `java.util.concurrent.atomic` — lock-free thread-safe programming on simple variables such as atomic integers, atomic booleans
 - ▶ `java.util.concurrent.locks` — framework for various locking algorithms, e.g., read -write locks and condition variables.

Thursday, August 27, 2009

Locks I

```
package java.util.concurrent.locks;
public interface Lock {
    public void lock(); // Wait for the lock to be acquired
    public Condition newCondition();
    // Create a new condition variable for use with the Lock
    public void unlock();
    ...
}

public class ReentrantLock implements Lock {
    public ReentrantLock();
    public void lock();
    public Condition newCondition();
    public void unlock();
}
```

Thursday, August 27, 2009

Locks II

```
package java.util.concurrent.locks;

public interface Condition {
    public void await() throws InterruptedException;
    //Atomically releases associated lock and cause thread to wait
    public void signal(); // Wake up one waiting thread
    public void signalAll(); // Wake up all waiting threads
}
```

Thursday, August 27, 2009

Generic Bounded Buffer I

```
class BoundedBuffer<Data> {

    private final Data buffer[];
    private int first, last, numberInBuffer;
    private final int size;
    private final Lock lock = new ReentrantLock();
    private final Condition notFull = lock.newCondition();
    private final Condition notEmpty = lock.newCondition();

    public BoundedBuffer(int length) {
        buffer = (Data[]) new Object[size = length];
    }
}
```

Thursday, August 27, 2009

Generic Bounded Buffer II

```

public void put(Data item) throws InterruptedException {
    lock.lock();
    try {
        while (numberInBuffer == size) notFull.await();
        last = (last + 1) % size;
        numberInBuffer++;
        buffer[last] = item;
        notEmpty.signal();
    } finally { lock.unlock(); }
}

public Data get() throws InterruptedException {
    lock.lock();
    try {
        while (numberInBuffer == 0) notEmpty.await();
        first = (first + 1) % size ;
        numberInBuffer--;
        notFull.signal();
        return buffer[first];
    } finally { lock.unlock(); }
}

```

Thursday, August 27, 2009

Asynchronous Thread Control

- Early versions of Java allowed one thread to asynchronously effect another thread through

```

public class Thread {
    ...
    public final void suspend();
    public final void resume();
    public final void stop();
    public final void stop(Throwable except)
        throws SecurityException;
}

```

All of the above methods are now obsolete
and therefore should not be used

Thursday, August 27, 2009

Thread Interruption

```
public class Thread ...  
    public void interrupt();  
        // Send an interrupt to the associated thread  
    public boolean isInterrupted();  
        // Returns true if associated thread has been  
        // interrupted, interrupt status is left unchanged  
  
    public static boolean interrupted();  
        // Returns true if the current thread has been  
        // interrupted and clears the interrupt status
```

Thursday, August 27, 2009

Thread Interruption

When a thread interrupts another thread:

- If the interrupted thread is blocked in wait, sleep or join, it is made runnable and the InterruptedException is thrown
- If the interrupted thread is executing, a flag is set indicating that an interrupt is outstanding; **there is no immediate effect on the interrupted thread**
- Instead, the called thread must periodically test to see if it has been interrupted using the isInterrupted or interrupted methods
 - ▶ If the thread doesn't test but attempts to blocks, it is made runnable immediately and the InterruptedException is thrown

Thursday, August 27, 2009

Summary

- True monitor condition variables are not directly supported by the language and have to be programmed explicitly
- Communication via unprotected data is inherently unsafe
- Asynchronous thread control allows thread to affect the progress of another without the threads agreeing in advance as to when that interaction will occur
- There are two aspects to this: suspend and resuming a thread (or stopping it all together), and interrupting a thread
- The former are now deemed to be unsafe due to their potential to cause deadlock and race conditions
- The latter is not responsive enough for real-time systems

Thursday, August 27, 2009

Completing The Java Model

- Aims:
 - To introduce thread priorities and thread scheduling
 - To show how threads delay themselves
 - To summarises the strengths and weaknesses of Java model
 - To introduce Bloch's safety levels

Thursday, August 27, 2009

Thread Priorities

- Although priorities can be given to Java threads, they are only used as a guide to the underlying scheduler when allocating resources
- An application, once running, can explicitly give up the processor resource by calling the **yield** method, placing the thread to the back of the run queue for its priority level

```
public class Thread ...  
    public static final int MAX_PRIORITY    = 10;  
    public static final int MIN_PRIORITY    = 1;  
    public static final int NORM_PRIORITY   = 5;  
  
    public final int getPriority();  
    public final void setPriority(int newPriority);  
    public static void yield();
```

Thursday, August 27, 2009

Warning

- From a real-time perspective, Java's scheduling and priority models are weak; in particular:
 - ▶ no guarantee is given that the highest priority runnable thread is always executing
 - ▶ equal priority threads may or may not be time sliced
 - ▶ where native threads are used, different Java priorities may be mapped to the same operating system priority

Thursday, August 27, 2009

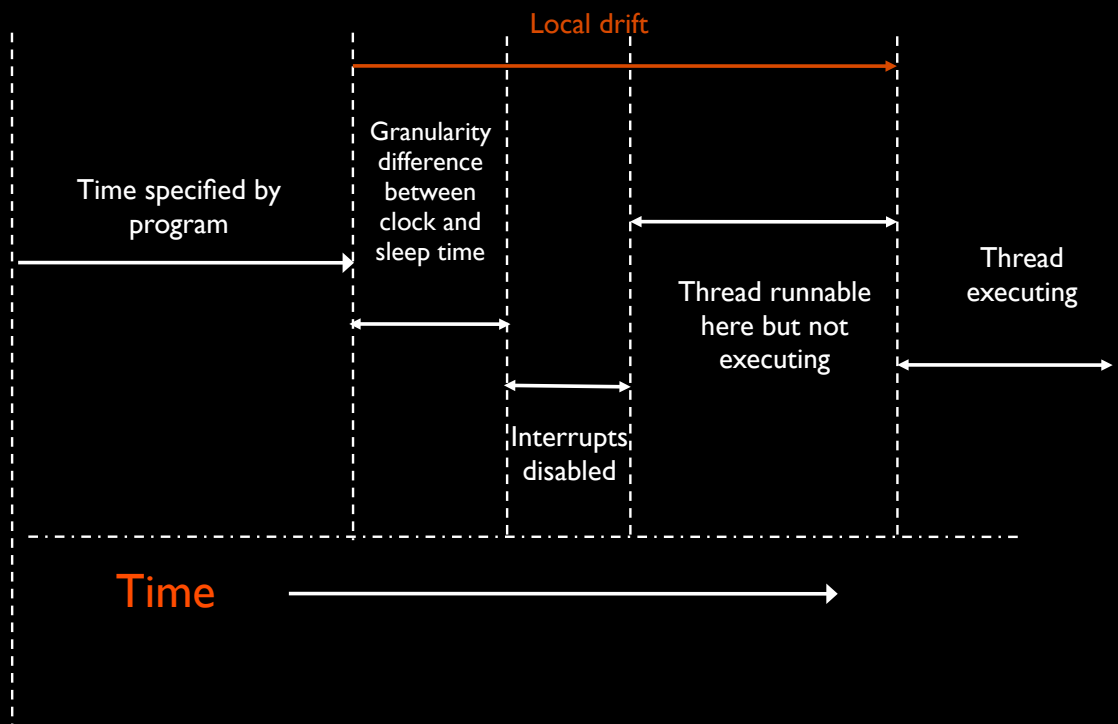
Delaying Threads: Clocks

- Java supports the notion of a wall clock
- `System.currentTimeMillis` returns the number of milliseconds since 1/1/1970 GMT and is used by `java.util.Date`
- However, a thread can only be delayed from executing by calling the `sleep` methods in the `Thread` class
- `sleep` provides a relative delay (sleep from now for some time), rather than sleep until 15th December 2003

```
class Thread... ..
    static void sleep(long ms) throws InterruptedException;
    static void sleep(long ms, int nanoseconds) throws Interrup
```

Thursday, August 27, 2009

Sleep Granularity



Thursday, August 27, 2009

Absolute Delays I

- Consider an embedded system where the software controller needs to invoke two actions
- The second action must occur a specified period (say 10 seconds) after the first action has been initiated
- Simply sleeping for 10 seconds after a call to the first action will not achieve the desired effect for two reasons
 - The first action may take some time to execute. If it took 1 second then a sleep of 10 would be a total delay of 11 seconds
 - The thread could be pre-empted after the first action and not execute again for several seconds
- This makes it extremely difficult to determine how long the relative delay should be

Thursday, August 27, 2009

Absolute Delays II

```
try{
    long start = System.currentTimeMillis();
    action_1();
    long end = System.currentTimeMillis();
    Thread.sleep(10000-(end-start));
} catch (InterruptedException ie) {...};
action_2();
```

What is wrong with this approach?

Thursday, August 27, 2009

Timeout on Waiting I

- In many situations, a thread can wait for an arbitrary long period time within synchronized code for an associated **notify**
- The absence of the call, within a specified period of time, sometimes requires that the thread take some alternative action
- Java provides two methods for this situation both of which allows the **wait** method call to timeout
- There are two important points to note
 - ▶ As with sleep, the timeout is a relative time and not an absolute time
 - ▶ It isn't possible to know if the thread is woken by timeout or notify

Thursday, August 27, 2009

Timeouts on Waiting

```
public class TimeoutException extends Exception {}

public class TimedWait {
    public static void wait(Object lock, long millis)
        throws InterruptedException, TimeoutException{
        // assumes the lock is held by the caller
        long start = System.currentTimeMillis();
        lock.wait(millis);
        if(System.currentTimeMillis() >= start + millis)
            throw new TimeoutException();
    }
}
```

What is wrong with this approach?

Thursday, August 27, 2009

Strengths of the Java Concurrency Model

- Main strength is simplicity and direct support by the language
- Many of errors that potentially occur with uses of an operating system interface for concurrency do not exists in Java
- Language syntax + strong type checking gives some protection e.g., it is not possible to forget to end a synchronized block
- Portability is enhanced as the concurrency model is the same irrespective of the operating system on which the program runs

Thursday, August 27, 2009

Weaknesses I

- Lack of support for condition variable
- Poor support for absolute time and time-outs on waiting
- No preference given to threads continuing after a notify over threads waiting to gain access to the monitor for the first time
- Poor support for priorities

Thursday, August 27, 2009

Weaknesses II

- Synchronized code should be kept as short as possible
- Nested monitor calls should be avoided because the outer lock is not released when the inner monitor waits; this can lead to deadlocks
- It is not always obvious when a nested monitor call is made:
 - non-synchronized methods can still contain a synchronized block
 - non-synchronized methods can be overridden with a synchronized method; method calls which start off unsynchronized may be used with a synchronized subclass
 - interface methods cannot be labelled as synchronized

Thursday, August 27, 2009

Bloch's Thread Safety Levels

- **Immutable** — Objects are constant and cannot be changed
- **Thread-safe** — Objects are mutable but they can be used safely in a concurrent environment as the methods are synchronized
- **Conditionally thread-safe** — Objects either have methods which are thread-safe, or have methods which are called in sequence with the lock held by the caller
- **Thread compatible** — Instances of the class provide no synchronization. However, instances of the class can be safely used in a concurrent environment, if the caller provides the synchronization by surrounding each method with the appropriate lock
- **Thread-hostile** — Instances should not be used in a concurrent environment even if the caller provides external synchronization. Typically because accessing static data or the external environment

Thursday, August 27, 2009