

# Atomicity for Reliable Concurrent Software

Cormac Flanagan  
UC Santa Cruz

Shaz Qadeer  
Microsoft Research

Joint work with  
Stephen Freund

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

1

## Towards Reliable Multithreaded Software

- Multithreaded software
  - increasingly common (Java, C#, GUIs, servers)
  - decrease latency
  - exploit underlying hardware
    - multi-core chips
- Heisenbugs due to thread interference
  - race conditions
  - atomicity violations
- Need tools to verify atomicity
  - dynamic analysis
  - type systems

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

2

## Motivations for Atomicity

### 1. Beyond Race Conditions

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

3

## Race Conditions

```
class Ref {  
  int i;  
  void inc() {  
    int t;  
    t = i;  
    i = t+1;  
  }  
}
```

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

4

## Race Conditions

```
class Ref {  
  int i;  
  void inc() {  
    int t;  
    t = i;  
    i = t+1;  
  }  
}  
  
Ref x = new Ref(0);  
  
x.inc();  
x.inc();  
  
assert x.i == 2;
```

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

5

## Race Conditions

```
class Ref {  
  int i;  
  void inc() {  
    int t;  
    t = i;  
    i = t+1;  
  }  
}  
  
Ref x = new Ref(0);  
parallel {  
  x.inc(); // two calls happen  
  x.inc(); // in parallel  
}  
assert x.i == 2;
```

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

6

- A **race condition** occurs if
- two threads access a shared variable at the same time
  - at least one of those accesses is a write

## Lock-Based Synchronization

```
class Ref {
  int i; // guarded by this
  void inc() {
    int t;
    synchronized (x) {
      t = i;
      i = t+1;
    }
  }
}

Ref x = new Ref(0);
parallel {
  x.inc(); // two calls happen
  x.inc(); // in parallel
}
assert x.i == 2;
```

- Field guarded by a lock
- Lock acquired before accessing field
- Ensures race freedom

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

7

## Limitations of Race-Freedom

```
class Ref {
  int i; // guarded by this
  void inc() {
    int t;
    synchronized (x) {
      t = i;
      i = t+1;
    }
  }
}

Ref x = new Ref(0);
parallel {
  x.inc(); // two calls happen
  x.inc(); // in parallel
}
assert x.i == 2;
```

- Ref.inc()
- race-free
  - behaves correctly in a multithreaded context

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

8

## Limitations of Race-Freedom

```
class Ref {
  int i;
  void inc() {
    int t;
    synchronized (this) {
      t = i;
    }
    synchronized (this) {
      i = t+1;
    }
  }
  ...
}
```

- Ref.inc()
- race-free
  - behaves **incorrectly** in a multithreaded context

Race freedom **does not** prevent errors due to unexpected interactions between threads

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

9

## Limitations of Race-Freedom

```
class Ref {
  int i;
  void inc() {
    int t;
    synchronized (this) {
      t = i;
      i = t+1;
    }
  }
  void read() { return i; }
  ...
}
```

- Ref.read()
- has a race condition
  - behaves **correctly** in a multithreaded context

Race freedom **is not necessary** to prevent errors due to unexpected interactions between threads

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

10

## Race-Freedom

- Race-freedom is neither *necessary* nor *sufficient* to ensure the absence of errors due to unexpected interactions between threads
- Is there a more fundamental semantic correctness property?

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

11

## Motivations for Atomicity

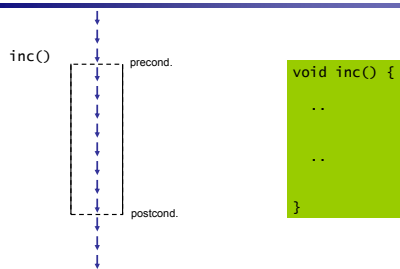
### 2. Enables Sequential Reasoning

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

12

## Sequential Program Execution




C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

13

## Multithreaded Execution




C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

14

## Multithreaded Execution

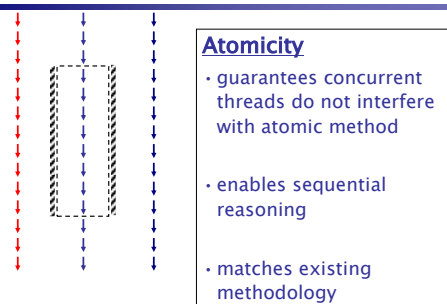


C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

15

## Multithreaded Execution



C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

16

## Motivations for Atomicity

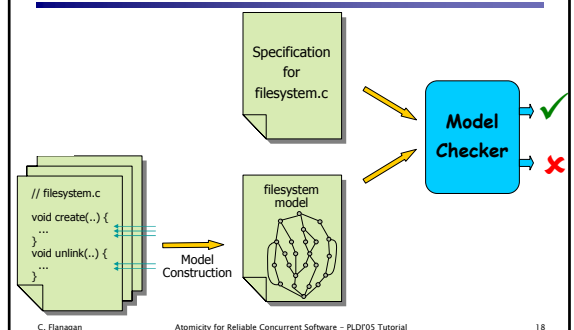
### 3. Simple Specification

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

17

## Model Checking of Software Models

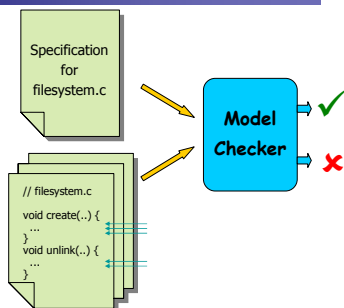


C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

18

## Model Checking of Software

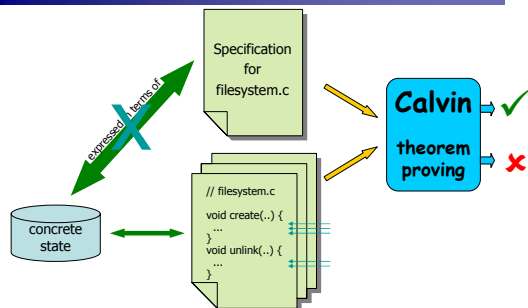


C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

19

## Experience with Calvin Software Checker

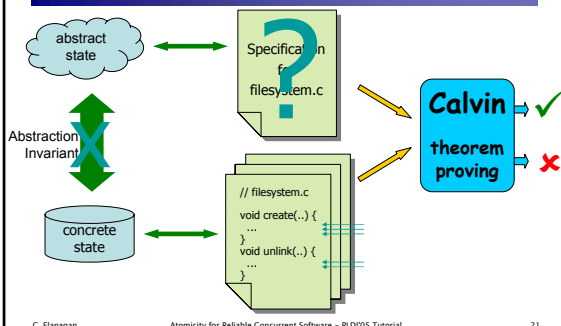


C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

20

## Experience with Calvin Software Checker




C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

21

## The Need for Atomicity

Sequential case:  
code inspection & testing mostly ok



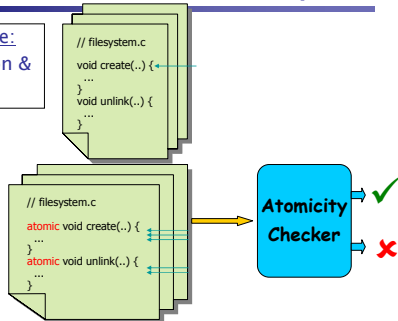
C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

22

## The Need for Atomicity

Sequential case:  
code inspection & testing ok



C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

23

## Motivations for Atomicity

1. Beyond Race Conditions
2. Enables Sequential Reasoning
3. Simple Specification

C. Flanagan


Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

24


## Atomicity

- The method `inc()` is **atomic** if concurrent threads do not interfere with its behavior

- Guarantees that for every execution



- there is a *serial* execution with same behavior



C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

25

## Atomicity


- Canonical property
  - (cmp. linearizability, serializability, ...)
- Enables sequential reasoning
  - simplifies validation of multithreaded code
- Matches practice in existing code
  - most methods (80%+) are atomic
  - many interfaces described as “thread-safe”
- Can verify atomicity statically or dynamically
  - atomicity violations often indicate errors
  - leverages Lipton’s theory of reduction

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

26

## Reduction [Lipton 75]



C. Flanagan


Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

27

## Checking Atomicity

```
atomic void inc() {
    int t;
    synchronized (this) {
        t = i;
        i = t + 1;
    }
}
```

R: right-mover    lock acquire  
L: left-mover    lock release  
B: both-mover    race-free variable access  
A: atomic        conflicting variable access



- Reducible blocks have form:  $(R|B)^* [A] (L|B)^*$

C. Flanagan


Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

28

## Checking Atomicity (cont.)

```
atomic void inc() {
    int t;
    synchronized (this) {
        t = i;
    }
    synchronized (this) {
        i = t + 1;
    }
}
```

R: right-mover    lock acquire  
L: left-mover    lock release  
B: both-mover    race-free variable access  
A: atomic        conflicting variable access



C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

29

## java.lang.StringBuffer

```
/**
 * ... used by the compiler to implement the binary
 * string concatenation operator ...
 */
```

String buffers are safe for use by multiple threads. The methods are synchronized so that all the operations on any particular instance behave as if they occur in some serial order that is consistent with the order of the method calls made by each of the individual threads involved.

```
/*
 * /* atomic */ public class StringBuffer { ... }
```

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

30

## java.lang.StringBuffer

```
public class StringBuffer {
    private int count;
    public synchronized int length() { return count; }
    public synchronized void getChars(...) { ... }

    atomic public synchronized void append(StringBuffer sb) {
        int len = sb.length(); // sb.length() acquires lock on sb,
        ...                    // gets length, and releases lock
        ...                    // other threads can change sb
        sb.getChars(..., len, ...);
        ...                    // use of stale len may yield
    }                          // StringIndexOutOfBoundsException
    }                          // inside getChars(...)
}
```

C. Flanagan

Atomicity for Reliable Concurrent Software – PLD/OS Tutorial

31

## java.lang.StringBuffer

```
public class StringBuffer {
    private int count;
    public synchronized int length() { return count; }
    public synchronized void getChars(...) { ... }

    atomic public synchronized void append(StringBuffer sb) {
        int len = sb.length();
        ...
        sb.getChars(..., len, ...);
        ...
    }
}
```

C. Flanagan

Atomicity for Reliable Concurrent Software – PLD/OS Tutorial

32

## Tutorial Outline

- Part 1
  - Introduction
  - Runtime analysis for atomicity
- Part 2
  - Model checking for atomicity
- Part 3
  - Type systems for concurrency and atomicity
- Part 4
  - Beyond reduction – atomicity via “purity”

C. Flanagan

Atomicity for Reliable Concurrent Software – PLD/OS Tutorial

33

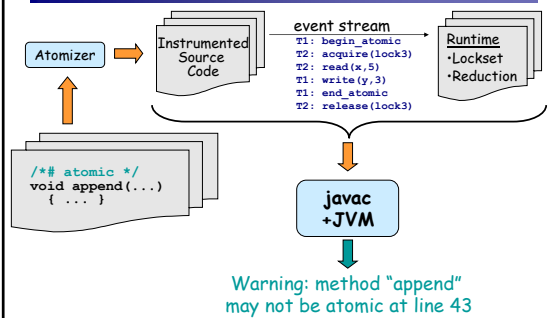
## Part I continued: Runtime Analysis for Atomicity

C. Flanagan

Atomicity for Reliable Concurrent Software – PLD/OS Tutorial

35

## Atomizer: Instrumentation Architecture



C. Flanagan

Atomicity for Reliable Concurrent Software – PLD/OS Tutorial

36

## Atomizer: Dynamic Analysis

- Lockset algorithm
  - from Eraser [Savage et al. 97]
  - identifies race conditions
- Reduction [Lipton 75]
  - proof technique for verifying atomicity, using information about race conditions

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

37

## Analysis 1: Lockset Algorithm

- Tracks *lockset* for each field
  - lockset = set of locks held on all accesses to field
- Dynamically infers protecting lock for each field
- Empty lockset indicates possible race condition

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

38

## Lockset Example

<p>Thread 1</p> <pre>synchronized(x) {   synchronized(y) {     o.f = 2;   }   o.f = 11; }</pre>	<p>Thread 2</p> <pre>synchronized(y) {   o.f = 2; }</pre>
---	---



- First access to  $o.f$ :

$$\text{LockSet}(o.f) = \text{Held}(\text{curThread})$$

$$= \{x, y\}$$

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

39

## Lockset Example

<p>Thread 1</p> <pre>synchronized(x) {   synchronized(y) {     o.f = 2;   }   o.f = 11; }</pre>	<p>Thread 2</p> <pre>synchronized(y) {   o.f = 2; }</pre>
---	---



- Subsequent access to  $o.f$ :

$$\text{LockSet}(o.f) := \text{LockSet}(o.f) \cap \text{Held}(\text{curThread})$$

$$= \{x, y\} \cap \{x\}$$

$$= \{x\}$$

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

40

## Lockset Example

<p>Thread 1</p> <pre>synchronized(x) {   synchronized(y) {     o.f = 2;   }   o.f = 11; }</pre>	<p>Thread 2</p> <pre>synchronized(y) {   o.f = 2; }</pre>
---	---



- Subsequent access to  $o.f$ :

$$\text{LockSet}(o.f) := \text{LockSet}(o.f) \cap \text{Held}(\text{curThread})$$

$$= \{x\} \cap \{y\}$$

$$= \{\}$$

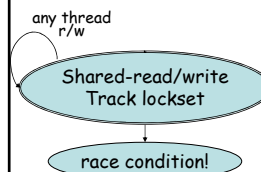
=> race condition

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

41

## Lockset

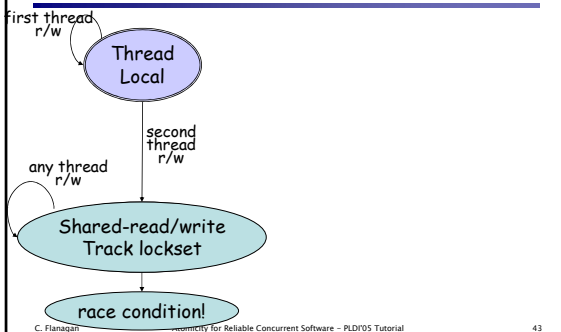


C. Flanagan

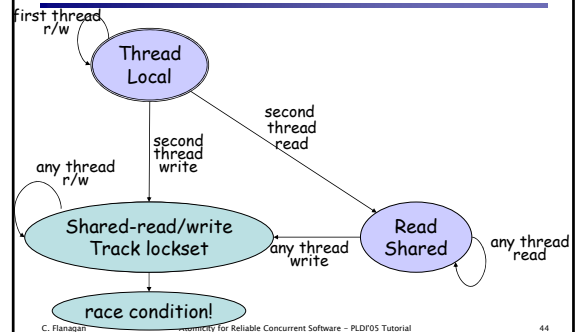
Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

42

## Lockset with Thread Local Data




## Lockset with Read Shared Data



## Atomizer: Dynamic Analysis


- Lockset algorithm
  - from Eraser [Savage et al. 97]
  - identifies race conditions
- Reduction [Lipton 75]
  - proof technique for verifying atomicity, using information about race conditions

## Reduction [Lipton 75]




## Performing Reduction Dynamically

- R: right-mover
  - lock acquire
- L: left-mover
  - lock release
- B: both-mover
  - race-free field access
- N: non-mover
  - access to "racy" fields



- Reducible methods:  $(R|B)^* [N] (L|B)^*$



## Atomizer Review

- Instrumented code calls Atomizer runtime
  - on field accesses, sync ops, etc
- Lockset algorithm identifies races
  - used to classify ops as movers or non-movers
- Atomizer checks reducibility of atomic blocks
  - warns about atomicity violations



## Evaluation

- 12 benchmarks
  - scientific computing, web server, std libraries, ...
  - 200,000+ lines of code
- Heuristics for atomicity
  - all synchronized blocks are atomic
  - all public methods are atomic, except `main` and `run`
- Slowdown: 1.5x – 40x

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

49

## Performance

Benchmark	Lines	Base Time (s)	Slowdown
elevator	500	11.2	-
hedc	29,900	6.4	-
tsp	700	1.9	21.8
sor	17,700	1.3	1.5
moldyn	1,300	90.6	1.5
montecarlo	3,600	6.4	2.7
raytracer	1,900	4.8	41.8
mrt	11,300	2.8	38.8
jigsaw	90,100	3.0	4.7
specJBB	30,500	26.2	12.1
webl	22,300	60.3	-
lib-java	75,305	96.5	-

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

50

## Extensions

- Redundant lock operations are both-movers
  - re-entrant acquire/release
  - operations on thread-local locks
  - operations on lock A, if lock B always acquired before A
- Write-protected data

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

51

## Write-Protected Data

```

class Account {
    int bal;
    /** atomic */ int read() { return bal; }
    /** atomic */ void deposit(int n) {
R      synchronized (this) {
B          int j = bal;
N          bal = j + n;
L      }
    }
}


```

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

52

## Extensions Reduce Number of Warnings



C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

53

## Evaluation

- Warnings: 97 (down from 341)
- Real errors (conservative): 7
- False alarms due to:
  - simplistic heuristics for atomicity
    - programmer should specify atomicity
  - false races
  - methods irreducible yet still "atomic"
    - eg caching, lazy initialization
- No warnings reported in more than 90% of exercised methods

C. Flanagan

Atomicity for Reliable Concurrent Software – PLDI'05 Tutorial

54

## java.lang.StringBuffer

```
public class StringBuffer {  
    private int count;  
    public synchronized int length() { return count; }  
    public synchronized void getChars(...) { ... }  
    /** atomic */  
    public synchronized void append(StringBuffer sb) {  
  
        int len = sb.length();  
        ...  
        ...  
        sb.getChars(..., len, ...);  
        ...  
    }  
}
```

StringBuffer.append is not atomic:  
Start:  
at StringBuffer.append(StringBuff  
at Thread1.run(Example.java:17)  
  
Commit: Lock Release  
at StringBuffer.length(StringBuff  
at StringBuffer.append(StringBuff  
at Thread1.run(Example.java:17)  
  
Error: Lock Acquire  
at StringBuffer.getChars(StringBu  
at StringBuffer.append(StringBuff  
at Thread1.run(Example.java:17)